



Technical Documentation

System Architecture, Development & Administration Guide

DOCUMENTATION SNAPSHOT · 2026-09-01

MEGA.EDU — Technical Documentation

Audience: developers, technical team members, system administrators, and future maintainers. **Status legend** (used throughout):  **Implemented** ·  **Designed/approved, not yet implemented** ·  **Known gap/issue** ·  **Future/planned** **Last verified:** 2026-09-01 (School Logos & User Profile Photos identity system), against the current codebase and the audited [/docs](#) documentation set. **Source discipline:** every claim in this document is drawn from the existing, individually-audited documents in [/docs](#) (cross-referenced throughout) and, where a doc was ambiguous, from direct inspection of the implementation. Nothing here describes a planned or hypothetical feature as if it were built. Where something is designed but not implemented, or implemented but deliberately incomplete, that is stated explicitly.

Table of Contents

1. Introduction & Purpose of This Document	5
2. System Overview	6
3. System Data Flow	9
4. Module Dependency Map	10
5. Technology Stack	12
6. Project Structure	13
7. Database Architecture & Prisma Models	15
8. Authentication & Authorization	18
9. User Roles & Permissions	20
10. School & Multi-School Architecture	21
11. Academic Sessions, Grades, Sections & Subjects	22
12. Student Lifecycle & GradeHistory	24
13. Teacher Assignments & Authorization Scopes	26
14. Attendance System	27
15. Evaluations	28
16. Parent–Teacher Meetings	29
17. Unit Tests	30
18. Assessment Framework System	31
19. Assessment Results & Publishing Workflow	33
20. Grading & GPA Calculations	35
21. Report Card Architecture	36
22. Promotion & Grade Decisions	37
23. Certificate System	39
24. School Logos & User Profile Photos	40
25. Dashboard Architecture	43
26. API Architecture & Major Routes	44
27. Business Rules & Data Integrity Protections	46

28. Audit Trails	47
29. Security Model	49
30. Calculation Engines & Shared Libraries	50
31. Deployment & Setup Requirements	51
32. Environment Variables	53
33. Database Migration & Update Procedures	54
34. Development Workflow & Verification	55
35. Demo & Sample Environment	56
36. Maintenance Guidelines	57
37. Known Gaps & Deliberate Out-of-Scope Decisions	58
38. Appendix: Documentation Index	60

1. Introduction & Purpose of This Document

This document is a single technical reference for MEGA.EDU, consolidating the detailed, individually-verified documents already maintained under [/docs](#) (each covering one subsystem in depth) into one narrative a new developer, system administrator, or maintainer can read top to bottom. It does not replace those documents — it cross-references them throughout, and the reader is expected to follow those links for full implementation detail (exact request/response shapes, full verification evidence, field-by-field schema notes). This document is the map; the individual [/docs](#) files are the territory.

Everything below reflects the system **as implemented and verified**, as of the phases completed through **Phase 3D-2/3/4** (Assessment Results, Publishing, Report Cards) plus the subsequent **guided Assessment Wizard, Class Overview** enhancements, and the **School Logos & User Profile Photos** identity system. Nothing here is invented, and nothing planned is described as available today — see [§37](#) for what is deliberately out of scope versus genuinely missing.


2. System Overview

MEGA.EDU is a national education network connecting schools, teachers, students, parents, and education organizations under one identity system, **MEGA ID**. A school gets a public profile and a digital identity; teachers and students at that school get accounts a school administrator approves; organizations can publish free online courses (MEGA Academy) and post opportunities; parents can follow their children's academic progress. Everyone signs in with the same account type, regardless of role.

The platform's purpose is to give every school, teacher, and student a **verifiable digital presence**: a school's identity isn't tied to one administrator's login, a student's academic record and certificates follow *them* rather than staying locked inside one school's paperwork, and achievements are independently verifiable by anyone with the right link.

High-level architecture

MEGA.EDU is a **single Next.js 14 App Router application** — there is no separate backend service. Pages and API routes are colocated under `src/app`, both talking to the same Prisma client against one database.




► **Technical Mermaid Source (for future editing)**

(Source: [ARCHITECTURE.md](#))

There is no client-side global state manager. `router.refresh()` — re-running the server component and re-fetching its data — is the application's entire "revalidate after mutation" strategy, following a `fetch()` → check `res.ok` → `router.refresh()` mutation flow throughout.

3. System Data Flow

The diagram below traces the major academic and assessment data flows a real school follows during a session — from setup through to a published Report Card and a promotion decision. It complements, rather than duplicates, the [Module Dependency Map](#) below: this shows the *order things happen in*, the map shows *what depends on what structurally*.




► Technical Mermaid Source (for future editing)

Reading this diagram: setup (session, grades, sections, subjects, teachers, students) must happen before any daily operations can run, since attendance, evaluations, units, and assessment entry are all scoped to a `GradeHistory` / `TeacherAcademicAssignment` combination that must already exist. Assessment results and Report Cards are the only flow that terminates in something a Parent/Student ever sees directly — every other flow (attendance, evaluations, unit tests) is visible to Student/Parent as read-only progress data, not a publish-gated artifact. Promotion is the flow that closes the loop back to a new `GradeHistory` row for the next session.

4. Module Dependency Map

This map shows structural dependencies between the major Phase 3 modules — an arrow means "the module at the arrow's head cannot function without data or authorization state from the module at its tail."



► Technical Mermaid Source (for future editing)

Solid arrows are hard, load-bearing dependencies (the dependent module's routes or queries directly reference the depended-on module's rows). Dotted arrows are looser, informational relationships (e.g. a `ParentTeacherMeeting` may link to an `StudentEvaluation`, but neither requires the other to exist).

Key structural facts this map summarizes:

- GradeHistory (Promotion/ PROMO) underlies everything.** Attendance, Unit Tests, Evaluations, and Assessment Results all resolve a student's current placement through `GradeHistory`, not an independently-maintained roster.
- TeacherAcademicAssignment / ClassTeacherAssignment gate almost every teacher-facing write.** Attendance and general Evaluations require `requireClassTeacher()`; Unit Tests, Subject Evaluations, and Assessment Results marks entry all require `requireTeacherAssignment()` scoped to the exact subject.

- **Assessment Results is the only module both Report Card and Class Overview ranking depend on.** Both reuse `fetchAssessmentResults()` / `buildReportCard()` — see §20 and §30.
- **Promotion (`recordGradeDecision()`) has no dependency on Assessment Results at all** — a Promotion decision can be recorded with zero assessment data ever entered. The relationship only exists in the *display* direction: Class Overview reads both the roster (from `GradeHistory`) and the ranking (from Assessment Results) on the same page, but neither module's write path touches the other.

5. Technology Stack

Layer	Technology	Notes
Framework	Next.js 14 (App Router)	Single deployment, pages and API routes colocated
UI	React 18, TypeScript	No separate frontend framework or SPA build
Styling	Tailwind CSS	Utility-first, no CSS-in-JS library
ORM	Prisma 5.20	Single <code>PrismaClient</code> singleton (<code>src/lib/prisma.ts</code>)
Database (dev)	SQLite	Single file, <code>prisma/dev.db</code>
Database (production, planned)	PostgreSQL 🗄️	Never configured or tested — see §31
Authentication	NextAuth 4	Single Credentials provider, JWT session strategy
Password hashing	<code>bcryptjs</code>	<code>bcrypt.compare</code> / <code>bcrypt.hash(password, 10)</code>
Validation	<code>zod</code>	Used on registration routes
Dev tooling	TypeScript compiler, ESLint (<code>next lint</code>), Prisma Studio	No automated test framework installed — see §34

(Source: [PROJECT_OVERVIEW.md](#), [DEPLOYMENT.md](#))

No Prisma `enum`s are used anywhere in this schema — SQLite's Prisma connector doesn't support them, even unused ones. Every status/type/role field is a plain `String`, with valid values documented in a comment above the field. This is a project-wide, explicitly-approved rule (see §27), not an oversight, and it is intended to remain the convention even after a future move to PostgreSQL (which does support enums).

6. Project Structure

```
src/
  app/                # pages (App Router) + API routes, colocated
  api/                # route.ts handlers, grouped by resource
  dashboard/          # role-specific dashboard pages/components
    setup/            # Initial School Setup wizard (Phase 2)
    grades/           # Class Overview / Promotion rosters + Pending queue
    sessions/new/     # New Session rollover
    academics/        # Subject catalog, grade offerings, teacher assignments,
                      # [gradeSubjectId]/ – Teaching Plans, Units, Tests
    attendance/       # Daily attendance marking + correction
    evaluations/      # General Student Evaluations + meeting scheduling
    meetings/         # Cross-role Parent-Teacher Meetings management
    students/[studentId]/ # Staff-only Student Profile page
    profile/          # Self-service "My Profile" – avatar, identity info
    assessment-frameworks/ # Landing page, guided wizard, Advanced management
    assessment-results/ # Marks entry, publish workflow
    report-card/[studentId]/ # Live Report Card view
  admin/              # Platform Admin-only pages
  ...                 # public pages: schools, courses, opportunities, etc.
  components/         # shared client components (SiteHeader, DashboardHero,
                      # AcademicProgressPanel, MeetingActions, certificate/)
  lib/                # server-side helpers – see §30 (includes uploads.ts)
  types/              # ambient type augmentation (next-auth.d.ts)
  prisma/
    schema.prisma     # single source of truth for the data model
    seed.ts           # idempotent bootstrap fixtures (upserts)
    seed-demo.ts      # optional, fuller demo/sample environment – see §35
```

(Source: [ARCHITECTURE.md](#))

Rendering pattern

Two consistent patterns are used throughout, never mixed within one page:

1. **Server component fetches Prisma directly, renders (or passes props to) a client component.**

The dominant pattern for a logged-in user's own dashboard data — `dashboard/page.tsx` branches on role, runs the relevant Prisma query, hands a fully-hydrated object to a `"use client"` component. Every Phase 2/3 configuration and roster page follows this.

2. **Client component fetches its own API route.** Used for mutations and independently-polled data (e.g. the notification bell). The mutation flow is consistently `fetch()` → check `res.ok` → `router.refresh()` .

The `src/lib` layer — "the only path" pattern

A recurring, deliberate architectural decision: for any write that must be atomic with a side effect, or must never be bypassed, the codebase centralizes it into one function, and every caller goes through it rather than calling `prisma.<model>.update()` directly. See §30 for the full inventory.

7. Database Architecture & Prisma Models

Datasource: SQLite in development (`prisma/dev.db`); the schema's own header comment and `.env.example` both mark PostgreSQL as the intended production target — nothing production-specific is configured yet. Every model below is  implemented (exists, migrated, and has at least one route reading/writing it) unless marked otherwise. Full field-by-field detail lives in [DATABASE.md](#); this section is an organized summary.

MEGA ID & roles

- `User` — the single identity record (email unique, `passwordHash`, `name`, `avatarUrl`). Every other model traces back to a `User` for who did what. `avatarUrl` lives here rather than on `Teacher` / `Student` / `Parent` specifically because a photo belongs to the person, not any one role — see §24.
- `UserRole` — which role(s) a user holds (`PLATFORM_ADMIN` | `SCHOOL_ADMIN` | `TEACHER` | `STUDENT` | `PARENT` | `ORGANIZATION_ADMIN` | `ACCOUNTANT`), `@@unique([userId, role])` — a user can hold several different roles at once.

Schools & Organizations

- `School` — `verified` / `isActive` flags ( `isActive` is read in two places but nothing ever sets it `false` — no deactivation action exists), `logoUrl` (uploadable/manageable by an authorized School Admin — see §24).
- `SchoolAdmin` / `SchoolAccountant` — join tables granting access to one school.
- `Organization`, `OrganizationAdmin` / `OrganizationAccountant` — the parallel structure for MEGA Academy publishers.  `Organization` has no `logoUrl` field in the schema at all.

People

- `Teacher`, `Student` — optional 1:1 with `User`, optional `schoolId`, `approved` flag. `Student.gradeLevel` is the **permanent legacy free-text grade fallback** — retained forever, never dropped, no longer written once a school completes Initial Setup.
- `Parent`, `ParentStudent` — a parent's linked children, matched by the child's email.

Academic Sessions & Grades (Phase 2)

- `AcademicSession` — one school-year window; at most one `ACTIVE` per school (app-enforced).
- `GradeReference` — the platform-wide, fixed 13-row grade ladder (`PP1` – `PP3`, `Y1` – `Y10`); not school-editable.
- `SchoolGrade` — a school's opt-in to one `GradeReference`, with its own `displayName`.

- **Section** — an optional, school-defined subdivision of a **SchoolGrade**; not session-scoped; soft-deactivate only, no hard delete.
- **TeacherGradeAssignment** — per-session teacher-to-grade link.
- **GradeHistory** — a student's grade placement for one session, the permanent record; **status:** **ENROLLED** | **COMPLETED** | **REPEATED** | **TRANSFERRED** | **LEFT**. See §12.
- **GradeHistoryAudit** — append-only log of every decision and section change ever written to a **GradeHistory** row.

Subjects & Teacher Academic Assignment (Phase 3A)

- **Subject** — school-wide catalog, reusable across every grade/session; deactivate only.
- **GradeSubject** — which subjects a grade offers, **for one specific academic session** — never carried forward.
- **TeacherAcademicAssignment** — teacher → session → grade → optional section → subject; a real **DELETE** route exists (current-state, not historical).

School Academic Operations (Phase 3B)

- **ClassTeacherAssignment** — Grade Class Teacher (**sectionId: null**) or Section Teacher; grade-wide and section-specific may coexist (unlike **TeacherAcademicAssignment**).
- **Attendance** / **AttendanceAudit** — one row per student per calendar day; audited corrections. See §14.
- **TeachingPlan** / **TeachingUnit** — a declared planned-total (separate model, not inferred from unit count) and the actual curriculum units with a teaching-progress status.
- **UnitTest** / **UnitTestResult** — pre-created roster pattern. See §17.

Teacher Qualitative Evaluation & Parent-Teacher Meetings (Phase 3C)

- **StudentEvaluation** / **StudentEvaluationAudit** — General vs. Subject via one nullable field; audit-on-share. See §15.
- **ParentTeacherMeeting** — periodic and occasional through one model. See §16.

Assessment Framework Foundation (Phase 3D-1)

- **AssessmentFramework** , **AssessmentPeriod** , **AssessmentComponent** , **GradingScale** , **GradingScaleBand** , **AssessmentFrameworkAssignment** . See §18.

Assessment Results, Publishing, Report Cards (Phase 3D-2/3/4)

- `AssessmentComponentResult`, `AssessmentComponentResultAudit`, `AssessmentResultPublication`, plus `GradingScaleBand.isPassing`. See §19.

MEGA Academy & Certificates

- `Course` → `CourseModule` → `Lesson`; `CourseEnrollment`; `Certificate`; `Instructor` (deliberately decoupled from `User` — can be named with no MEGA ID). See §23.

Identity layer, content, and other supporting models

- `Interest` (self-declared), `Skill` (teacher-attested, `@@unique([studentId, addedByUserId, name])`).
- `Program`, `NewsPost`, `Resource`, `Event`, `Opportunity` — simple content types, no approval workflow.
- `Subscription` / `Payment` 🟡 — fully modeled, **no payment processor is connected**; priced-course enrollment is explicitly blocked.
- `Notification` — written only through `notify()` / `notifySchoolCommunity()`, never directly.

Cross-cutting schema conventions

- **Snapshot fields for anything that must survive a later rename:** `Certificate.*NameSnapshot`, `GradeHistoryAudit.previous*/new*` are plain values, not live FK relations. Logos are the one deliberate exception (looked up live).
- **The nullable-scope-discriminator idiom**, used identically across five different models: `TeacherAcademicAssignment.sectionId`, `TeachingUnit.sectionId`, `StudentEvaluation.gradeSubjectId`, `ParentTeacherMeeting.gradeSubjectId`, `AssessmentFrameworkAssignment.gradeSubjectId` — `null` means "general/grade-wide/default", a real value means "specific/override".
- **The recurring `NULL ≠ NULL` unique-index gap:** SQL treats two `NULL` values in a unique index as non-colliding, so a `@@unique` constraint with a nullable column cannot by itself block two rows that are both `null` in that column. This gap has been found (twice, live) and pre-checked (three times, proactively) across `TeacherAcademicAssignment`, `ClassTeacherAssignment`, `StudentEvaluation`, `AssessmentFrameworkAssignment`, and `AssessmentComponent` — each has an explicit application-level `findFirst`-based pre-check function rather than relying on the database constraint alone.
- **Most models have no working delete route** — for those, "delete behavior" is schema-level cascade behavior that would apply if a delete ever happened, never something a UI action triggers. The deliberate exceptions (real `DELETE` routes, because they're current-state operational data, not historical record): `TeacherGradeAssignment`, `GradeSubject`, `TeacherAcademicAssignment`, `ClassTeacherAssignment`, `AssessmentPeriod`, `AssessmentComponent` (until results exist), `AssessmentFrameworkAssignment` (until results exist).

(Source: [DATABASE.md](#))

8. Authentication & Authorization

Login

`src/lib/auth.ts` — NextAuth v4 with a single `CredentialsProvider` ("MEGA ID"): email + password, `bcrypt.compare` against `User.passwordHash`. `PrismaAdapter` is attached, but the session strategy is **JWT**, not database sessions. Custom sign-in page at `/login`.

Session handling

Every server-side check reads the session identically:

```
const session = await getServerSession(authOptions);
const userId = (session?.user as any)?.id;
const roles = (session?.user as any)?.roles;
```

The `jwt` callback copies `id` / `roles` onto the token at sign-in; the `session` callback copies them back onto `session.user`.

The `requireX()` helper suite (`src/lib/authorize.ts`)

Every write route that needs "does this session have permission to act on this specific resource" calls one of these rather than inlining the check. Each returns the authorized `userId` or `null` (never throws).

Function	Checks	Used for
<code>requireSchoolAdmin(schoolId)</code>	Session + a <code>SchoolAdmin</code> row	Nearly every <code>/api/schools/[id]/*</code> write route
<code>requireOrgAdmin(organizationId)</code>	Session + an <code>OrganizationAdmin</code> row	Every <code>/api/organizations/[id]/*</code> write route
<code>requireCourseOwner(courseId)</code>	Resolves course → org → org-admin check	Course content routes
<code>requirePlatformAdmin()</code>	Session + <code>roles.includes("PLATFORM_ADMIN")</code>	<code>/admin/*</code> , <code>/api/admin/*</code>
<code>requireSchoolFinance(schoolId)</code>	Session + (<code>SchoolAdmin</code> or <code>SchoolAccountant</code>)	Finance-only surfaces

<code>requireOrgFinance(organizationId)</code>	Session + (<code>OrganizationAdmin</code> or <code>OrganizationAccountant</code>)	Finance-only surfaces
<code>requireTeacherAssignment(schoolId, scope)</code>	Approved <code>Teacher</code> + a matching <code>TeacherAcademicAssignment</code>	Teaching Plans/Units/Tests, Subject Evaluations, Assessment Results marks entry/publish
<code>requireClassTeacher(schoolId, scope)</code>	Approved <code>Teacher</code> + a matching <code>ClassTeacherAssignment</code>	Attendance, General Evaluations
<code>teacherHoldsSubjectAssignment()</code> / <code>teacherHoldsClassAssignment()</code>	A named teacher's assignment (not the session's own user)	Admin-attributed evaluation/meeting creation

Deliberate design note: the finance helpers check *both* the Admin and Accountant relationships on purpose — an Admin keeps full authority (finance included); a bare Accountant gets finance access *only*.

The three-way `sectionId` scope semantics

`requireTeacherAssignment()` / `requireClassTeacher()` both accept an optional `sectionId` scope with **three distinct states**, resolved via a shared `sectionScopeWhere()` helper:

- **omitted** — no section restriction, matches any assignment for the grade/subject.
- **null** — the target itself is grade-wide (e.g. a grade-wide `TeachingUnit`); requires a grade-wide assignment specifically.
- **a real section id** — the target is one specific section; a grade-wide assignment (covers every section) or that exact section's assignment both satisfy it.

This distinction was corrected in Phase 3B, before either helper had a real caller — the original implementation collapsed `null` and *omitted* into one code path (both are falsy in JavaScript), which would have wrongly authorized a section-specific-only teacher to manage a grade-wide unit.

Access patterns not covered by `authorize.ts`

Some checks are simple/specific enough to stay inlined: certificate preview access (`recipientUserId === userId || PLATFORM_ADMIN`), course enrollment/completion ownership, the Promotion roster's closed-session access, and the Student Profile page's staff-wide (not assignment-scoped) access check.

(Source: [AUTHENTICATION_AND_AUTHORIZATION.md](#))

9. User Roles & Permissions

All seven roles are stored identically as `UserRole` rows; a single `User` can hold **multiple roles at once**. `dashboard/page.tsx` picks which dashboard to render using a fixed priority order, explicitly commented in the code as *"a simple MVP priority order, not a permission hierarchy"*:

PLATFORM_ADMIN → SCHOOL_ADMIN → TEACHER → STUDENT → PARENT → ORGANIZATION_ADMIN → ACCOUNTANT

Role	Core capabilities	Approval
PLATFORM_ADMIN	Command-center dashboard (real counts, no invented metrics); verifies schools/organizations; can view any certificate. ⚠️ Can only be granted via <code>seed.ts</code> or direct DB access — no in-app route.	Seeded fixture only
SCHOOL_ADMIN	Full school management: profile, staff/student approval, Phase 2 (sessions/grades/sections/promotion), Phase 3A (subjects/teacher assignments), Phase 3B (attendance/units/tests), Phase 3C (evaluations/meetings), Phase 3D (assessment config/results), direct Student/Teacher creation.	Self-registers as first admin
TEACHER	Skill-crediting; course enrollment; within assignment scope — teaching plans/units/tests, evaluations, attendance (if also a Class/Section Teacher), assessment marks entry/publish.	School Admin (skipped if created via <code>+ Add Teacher</code>)
STUDENT	Read-only: Teaching Progress, Test Results, Recent Attendance, shared Evaluations, published Assessment Results, Report Card, course enrollment/certificates.	School Admin (skipped if created via <code>+ Add Student</code>)
PARENT	Read-only, per linked child: the same progress views as Student, plus Parent-Teacher Meetings (Student never sees these).	N/A (self-links children)
ORGANIZATION_ADMIN	Course authoring/publishing, opportunities, accountant management. ⚠️ <code>Organization.verified</code> is never checked before publishing or enrollment.	Platform Admin (<code>verified</code> flag)
ACCOUNTANT	Finance-surface visibility only (no real transaction data exists — payments aren't integrated). Not self-registerable.	Granted directly by an Admin

(Source: [USER_ROLES.md](#), [MEGA_ID.md](#))

10. School & Multi-School Architecture

A **MEGA ID belongs to the individual**, not to a school — a `User` is never owned by an institution. School affiliation is a state on a `Teacher` / `Student` profile (`schoolId` , unaffiliated/pending/approved), not baked into the account itself. This means:

- **Learning history follows the person, not the school.** `CourseEnrollment` and `GradeHistory` both link to `Student.id → User.id` ; if a student transfers or leaves, their historical rows are never deleted — only a new `status` (`TRANSFERRED` / `LEFT`) is recorded.
- **A school and an organization are structurally independent.** A `Course` belongs only to an `Organization` ; Phase 2's grade structure belongs entirely to a `School` . The two only meet at `Certificate.associatedSchoolId` , an informational link, never an ownership relation.
- **Registration is deliberately two-stage.** Every path creates a minimal, single-role account first; school/organization affiliation happens afterward from the person's own dashboard (`POST /api/teacher/join-school` , `.../student/join-school` , `.../parent/link-child` , `.../schools/create-for-admin`).
- **Multi-school independence:** nothing in the schema assumes a single school's data ever interacts with another school's. `AcademicSession` , `SchoolGrade` , `Section` , `Subject` are all `schoolId` -scoped, and every write route resolves the caller's own school via `requireSchoolAdmin` / `requireTeacherAssignment` / `requireClassTeacher` before touching any row.

(Source: [MEGA_ID.md](#), [ARCHITECTURE.md](#))

11. Academic Sessions, Grades, Sections & Subjects

Academic Sessions

An `AcademicSession` represents one school year for one school — the container everything else in the academic system is scoped to. A school may have **at most one ACTIVE session** at a time, enforced at the application level (SQLite cannot express a partial unique index). Creating a session when one is already active returns `alreadyActive: true` (HTTP 200, not an error) rather than a second row — verified live against a genuine two-tab race condition.

Rollover (`POST /api/schools/[id]/academic-sessions/rollover` , `/dashboard/sessions/new`) closes the current session and opens a new one in one transaction: every student whose most recent decision was `COMPLETED` / `REPEATED` *with a real outcome grade* is auto-placed into the new session (`carryForwardEligibleStudents()`); anyone still `ENROLLED` with no decision is left unplaced and surfaces in the persistent Pending/Unresolved queue. **Section is never carried forward** — the new row's `sectionId` is always `null` , regardless of the student's prior section.

Grade decisions (Promotion) happen **independently of session status** — a School Admin can promote Grade 6 today and leave Grade 7 untouched for weeks; nothing forces every grade to be resolved before the next rollover.

Grades

`SchoolGrade` is a school's opt-in to one platform-wide `GradeReference` (the fixed 13-rung ladder `PP1 – PP3` , `Y1 – Y10`), with its own `displayName` (e.g. `"Class 6"`).

Sections

An **optional** subdivision of a `SchoolGrade` (Class 6 → A, B, C), added after the original Phase 2 build, on the same audit architecture:

- Belongs to the `SchoolGrade` , not to a session — the same rows persist across every academic session that grade is used in.
- Unlimited sections per grade; unique name per grade; **no hard-delete path** — deactivate (`isActive: false`) only, so historical `GradeHistory` / `GradeHistoryAudit` references always stay resolvable.
- School-Admin-only for every section-related action.
- **Never inherited across sessions or decisions** — a promotion, a repeat, and the rollover carry-forward sweep all create/leave the row's `sectionId` untouched from their own perspective (a decision never touches it) or explicitly `null` (a new session's row never inherits it). Section is always a separate, explicit action per session, mirroring how teacher assignments are never carried forward either.

Subjects (Phase 3A)

- **Subject** — a school-wide catalog entry, reusable across every grade and session; deactivate only, no hard delete.
- **GradeSubject** — which subjects a grade offers, **for one specific academic session** — every session starts with zero offerings for every grade; nothing is auto-copied from the prior session, so a past session's curriculum stays permanently reconstructable. A real **DELETE** route exists (blocked with **409** if a **TeacherAcademicAssignment** still references it).

(Source: [ACADEMIC_SESSIONS.md](#), [GRADES_AND_PROMOTION.md](#), [ACADEMIC_STRUCTURE.md](#))

12. Student Lifecycle & GradeHistory

`GradeHistory` is the central, permanent record of a student's grade placement for one session (`@@unique([studentId, academicSessionId])`). It is never deleted, and its `status` / `outcomeGradeId` may **only ever** be written through `recordGradeDecision()` (`src/lib/gradeHistory.ts`) — the sole audited write path. `sectionId` on an *existing* row may only ever be written through the separate `reassignSection()` function.

Lifecycle states

```
(no row) --create--> ENROLLED --decision--> COMPLETED | REPEATED | TRANSFERRED | LEFT
```

- **ENROLLED** — the default at creation; the only status eligible for a *new* Promotion decision.
- **COMPLETED** — promoted; `outcomeGradeId` required (the target grade).
- **REPEATED** — staying in the same grade next session; `outcomeGradeId` required (same grade).
- **TRANSFERRED** — left for another school; no outcome grade.
- **LEFT** — left the school entirely; no outcome grade.

First-time placement — four entry points, one shared architecture

A student's *first* `GradeHistory` row for a session can be created four ways, all sharing the identical "creation isn't a decision" rule (direct `GradeHistory.create()`, `status: "ENROLLED"`, no `decidedAt`, never routed through `recordGradeDecision()`):

1. **Initial Setup step 6** — `matchLegacyGradeText()` confident-match + manual queue, bulk, via `POST /api/schools/[id]/grade-placements`.
2. **Pending/Unresolved queue's "manually place"** — bulk, same route reused.
3. **Add Student, optionally, at creation time** — an inline `GradeHistory.create()` inside `POST /api/schools/[id]/students`'s own transaction.
4. **Students tab's "Assign Grade & Section →"** — a single-item call to the same `grade-placements` route, for any approved student with no `GradeHistory` row in the active session.

`CURRENT_ROSTER_STATUSES` — who counts as "currently in a grade"

`CURRENT_ROSTER_STATUSES` (`src/lib/gradeHistory.ts`) = `["ENROLLED", "COMPLETED", "REPEATED"]` — the shared definition of "physically in this grade for the rest of this session," used identically by the Grades index's "N enrolled" count and the Class Overview roster. `COMPLETED` / `REPEATED` are included

because a decision about *next* session doesn't remove a student from *this* grade the moment it's recorded — the school year isn't over. `TRANSFERRED` / `LEFT` remain excluded, since those students have genuinely left. This broadened definition **does not** change Promotion-panel eligibility — see [§22](#).

The legacy matching utility

`matchLegacyGradeText()` (`src/lib/gradeMatching.ts`) matches free text (`"Grade 6"` , `"VI"` , `"Nursery"`) to a `GradeReference` code, handling the full Roman-numeral range with subtractive notation, ordinal suffixes, and pre-primary keywords. **Returns** `null` — **never a guess — for anything ambiguous**, verified with 20+ real inputs.

(Source: [GRADES_AND_PROMOTION.md](#), [DATABASE.md](#))

13. Teacher Assignments & Authorization Scopes

Three distinct assignment models, each with a different overlap policy:

Model	Scope	Overlap policy	Carried forward?
<code>TeacherGradeAssignment</code>	teacher → grade → session	N/A (no subject/session)	No
<code>TeacherAcademicAssignment</code>	teacher → grade → optional section → subject → session	Exclusive: a teacher may not hold both a grade-wide and section-specific row for the same subject/grade/session	No
<code>ClassTeacherAssignment</code>	teacher → grade → optional section → session (Grade Class Teacher / Section Teacher)	Coexisting: grade-wide and section-specific rows may both exist for the same grade	No

`TeacherAcademicAssignment.gradeSubjectId` is a direct FK to the matching `GradeSubject` row — it is schema-impossible to assign a teacher to a subject the grade doesn't actually offer that session. Multiple different teachers may freely overlap on the same subject/grade/session (no teaching hierarchy exists or is planned).

`requireTeacherAssignment()` was built in Phase 3A *ahead of its first caller*, specifically as shared foundation for later phases — Phase 3B's Teaching Plans/Units/Tests, Phase 3C's Subject Evaluations, and Phase 3D-2's marks entry/publish all depend on it. See [§8](#) for its exact scope semantics.

(Source: [ACADEMIC_STRUCTURE.md](#), [ACADEMIC_OPERATIONS.md](#))

14. Attendance System

Attendance — one row per student **per calendar day**, never subject-based, `@@unique([studentId, date])` (global, not per-session). `date` is always derived from an explicit client-sent `"YYYY-MM-DD"` string, never the server's own clock, so the value represents the school's intended calendar day regardless of server timezone.

Marking: `POST /api/schools/[id]/attendance`, `requireSchoolAdmin` OR `requireClassTeacher` scoped to `sectionId` (omitted = whole grade, requires a Grade Class Teacher). A student whose actual placement doesn't match the marking pass's target grade/section is silently skipped; an already-marked student for that date is skipped, not an error.

Corrections go only through `correctAttendance()` (`src/lib/attendance.ts`), which updates the row **and** inserts an `AttendanceAudit` row capturing `previousStatus` / `newStatus` / `previousRemarks` / `newRemarks` together — every time, even a remarks-only edit records status unchanged, and vice versa.

Visibility: Student sees their own last-15-days attendance (read-only); Parent sees the same, per linked child; both reflect corrections, never the original pre-correction value.

(Source: [ACADEMIC_OPERATIONS.md](#))

15. Evaluations

`StudentEvaluation` — a teacher's narrative, qualitative remark about a student, for one session.

General vs. Subject is not a type field — it is entirely the presence or absence of `gradeSubjectId`:

- `gradeSubjectId: null` — **General Student Evaluation**, authored by a Grade Class Teacher / Section Teacher (`requireClassTeacher()`), on `/dashboard/evaluations` .
- `gradeSubjectId: <id>` — **Subject Evaluation**, authored by a Subject Teacher (`requireTeacherAssignment()`), on `/dashboard/academics/[gradeSubjectId]` .

A School Admin may create either **on behalf of a named teacher** — the named `teacherId` is independently validated (`teacherHoldsSubjectAssignment()` / `teacherHoldsClassAssignment()`), never simply trusted because the caller is an admin.

Audit-on-share

`updateEvaluationRemarks()` is the only code path that changes an existing evaluation's `remarks` :

- **While fully private** (`visibleToParent` and `visibleToStudent` both `false`) — a plain update, no audit row.
- **Once shared with EITHER audience** — every subsequent edit updates the row **and** inserts a `StudentEvaluationAudit` row, in the same transaction.

`visibleToParent` / `visibleToStudent` are **two fully independent gates** — sharing with one never implies the other; each has its own timestamp. Sharing itself is **one-way** — no un-share path exists in this phase.

The "STAFF" audience

Added for the Student Profile page (`/dashboard/students/[studentId]`) — skips the visibility filter entirely, returning every evaluation including private ones, since School Admins/approved Teachers already have full write access at their school.

(Source: [ASSESSMENT_AND_EVALUATION.md](#))

16. Parent–Teacher Meetings

`ParentTeacherMeeting` handles both **occasional** (single) and **periodic** (bulk) meetings through one model and one route (`POST /api/schools/[id]/meetings` , `{meetings: [...]}` — one item vs. many). No recurrence/series concept exists.

- **Initiation**: School Admin or an authorized Teacher only. **Parents have no write path at all** — read-only recipients in this phase.
- **Every item is resolved/validated before the transaction opens** — the Postgres-safe pattern, not the SQLite-only catch-mid-transaction pattern used by `grade-placements` / `teacher-assignments` .
- **Editing** (`PATCH .../meetings/[meetingId]`) is authorized by identity — the meeting's own (still-approved) teacher, or a School Admin.
- **Rescheduling** (`scheduledAt` / `location` / `onlineUrl`) is only allowed while `status: "SCHEDULED"` — a `COMPLETED` / `CANCELLED` meeting's details are historical record.
- `linkedEvaluationId` — a plain, non-unique FK to a `StudentEvaluation` , for attaching prepared context.

Structural, not just UI, visibility guarantee

`fetchMeetingsForStudent(studentId, audience)` (`src/lib/academicProgress.ts`) accepts **only** `"PARENT"` | `"STAFF"` — its type signature has no `"STUDENT"` member at all. It is called only from the Parent dashboard branch and the staff-only Student Profile page — never from the Student branch, and never folded into the shared `AcademicProgressPanel.tsx` that the Student's own render path uses. A Student attempting direct navigation to `/dashboard/students/[studentId]` or `/dashboard/meetings` is redirected away before any meeting data is fetched.

(Source: [ASSESSMENT_AND_EVALUATION.md](#))

17. Unit Tests

TeachingUnit — one curriculum unit/chapter under a **GradeSubject**, with **status: NOT_STARTED | IN_PROGRESS | COMPLETED** (managing **startedAt** / **completedAt** automatically) and an app-assigned **order**. **sectionId: null** means a grade-wide unit sequence shared by every section; a real value means that section's own independent sequence.

TeachingPlan — a **separate model**, not a field on **TeachingUnit**: a teacher's declared **plannedTotal** (independent of how many **TeachingUnit** rows actually exist) and **unitLabel** (free text — "Unit" or "Chapter", a plain string, not a second model).

UnitTest — belongs to one **TeachingUnit**; **creatable only once the unit is IN_PROGRESS or COMPLETED**. **UnitTestResult** is **pre-created** (**status: "PENDING"**) for every student currently enrolled in the test's scope at test-creation time — a stable roster snapshot, deliberately the opposite pattern from **AssessmentComponentResult**'s lazy creation (see §19) because a test has one clean, one-shot creation moment. **ABSENT** forces **marksObtained** to **null**; **EVALUATED** requires $0 \leq \text{marksObtained} \leq \text{maxMarks}$.

UnitTest / **UnitTestResult** remain **genuinely separate from AssessmentFramework / AssessmentComponent** — no merge, no auto-derivation of marks between the two systems, by explicit design decision.

(Source: [ACADEMIC_OPERATIONS.md](#), [ASSESSMENT_FRAMEWORK.md](#))

18. Assessment Framework System

Real report cards from different schools use genuinely different marking structures — percentage-weighted with letter grades, raw maximum marks, term-split, or a different structure per subject. Phase 3D-1 built a **configurable** assessment system rather than one hard-coded structure.

Core models

- **AssessmentFramework** — a reusable, **school-wide** marking-scheme template, **not session-scoped** (like **GradingScale**, and like **Subject** / **GradeSubject** 's catalog-vs-offering split). Deactivate only, no hard delete.
- **AssessmentPeriod** — an optional grouping layer ("Term I", "Annual") — a framework may define zero periods (flat structure) or several.
- **AssessmentComponent** — one scored/graded/descriptive piece, belonging directly to a framework or to one of its periods. **maxMarks: Float** does **double duty as marks and weight** — a component worth 10% and a component worth 10 raw marks are the identical shape; only **entryMode** (**MARKS** | **GRADE** | **DESCRIPTIVE**) differs.
- **GradingScale** / **GradingScaleBand** — a reusable, school-wide marks→grade conversion table (**minPercent** / **maxPercent** / **label** / **gradePoint?** / **isPassing?**).
- **AssessmentFrameworkAssignment** — the **only session-scoped model** in this phase: binds a framework to (**AcademicSession**, **SchoolGrade**) , optionally narrowed to one **GradeSubject** as a subject override.

Resolution rule

resolveFrameworkAssignment() (**src/lib/assessmentFramework.ts**): look up a subject-specific assignment first; if none exists, fall back to the grade-default. The same nullable-scope-discriminator idiom used four other places in this schema.

Setup UI — guided wizard as the primary path

/dashboard/assessment-frameworks/new — **"Create Assessment System"** is the primary, promoted way a School Admin sets up a framework: a 6-step guided flow in plain language (Name it → Terms or one overall assessment? → How are students assessed? → How should results be shown? → Where should this apply? → Review & confirm), deliberately hiding the underlying "framework"/"assignment" terminology. **Zero new API routes** — confirming the wizard sequences the exact same existing calls: an optional **POST /grading-scales**, then **POST /assessment-frameworks**, then **POST /assessment-framework-assignments**.

</dashboard/assessment-frameworks> itself is the landing page: a plain-language "Your Assessment Systems" list with the wizard promoted above it, and the original detailed editor (grading scales, per-framework period/component editing, the raw assignment form) tucked behind a collapsed **"Advanced management"** disclosure for anyone editing something after creation.

Authorization

School-Admin-only for every write route in this phase — no teacher-facing action exists until marks entry (Phase 3D-2).

(Source: [ASSESSMENT_FRAMEWORK.md](#))

19. Assessment Results & Publishing Workflow

Lazy result creation

Unlike `UnitTestResult`'s eager pre-creation, `AssessmentComponentResult` rows are created **only when a value is actually written**. A component is defined at framework-*design* time — long before any assignment exists — and can be added to an already-assigned, reusable framework; eager pre-creation would silently miss any component added afterward. The marks-entry page instead computes a **virtual roster** (`GradeHistory`, scoped to the assignment's grade) and left-joins whatever result rows already exist.

Publishing workflow

(no row) → first entry → DRAFT → explicit Publish → PUBLISHED → correction → PUBLISHED + audit

- **Publication is subject-level** (`AssessmentResultPublication`, keyed by `gradeSubjectId` + `studentId`), **never per-component** — a Parent/Student can never see a partially-published subject.
- **Who can publish:** `requireSchoolAdmin(id) || requireTeacherAssignment(id, {..., subjectId})` — the identical composition already used for `StudentEvaluation` sharing. Being a Class Teacher or Section Teacher alone does not provide special authority to enter, correct, or publish assessment results — assessment authority requires the appropriate matching `TeacherAcademicAssignment` for the subject.
- **Completeness is computed server-side**, not trusted from the client: `POST .../subjects/[gradeSubjectId]/publish` silently skips (never errors) any student with a still-`PENDING` non-`DESCRIPTIVE` component.
- **A correction to an already-`PUBLISHED` result stays published and is audited** (`AssessmentComponentResultAudit`) — it is **never reverted to `DRAFT`**, matching the identical `StudentEvaluation` precedent.

ABSENT handling

`status: "ABSENT"` is settable regardless of `entryMode`, forces `marksObtained` / `gradeLabel` / `remarks` to `null` regardless of what's passed, and contributes **zero** (not `null`) to numeric aggregation.

Data Integrity Protections (three guards on existing 3D-1 routes)

Guard	Rule
Component structural lock	<code>maxMarks</code> / <code>entryMode</code> locked once any result exists; renaming stays free

Assignment deletion block	<code>DELETE</code> rejected once any result/publication references it
Grading-scale lock	Full <code>bands</code> replacement rejected once any <code>PUBLISHED</code> result uses the scale; <code>name</code> / <code>isActive</code> remain editable

(Source: [ASSESSMENT_RESULTS.md](#))

20. Grading & GPA Calculations

All calculation lives in `src/lib/assessmentResults.ts` — **nothing is cached**; every total, percentage, and grade is recalculated on every read.

Function	Purpose
<code>computeComponentContribution()</code>	Turns one raw result into a point value. <code>MARKS</code> → raw value; <code>GRADE</code> → the matched <code>GradingScaleBand</code> 's percentage midpoint applied to <code>maxMarks</code> (deliberately not <code>gradePoint</code> , which is an arbitrary school-chosen scale with no inherent percentage relationship); <code>DESCRIPTIVE</code> → always excluded from every numeric total; <code>ABSENT</code> → zero; <code>PENDING</code> → <code>null</code> (never zero).
<code>aggregateGroup()</code>	One function for both period-scoped and framework-level aggregation — components → (periods) → subject, called identically at every level.
<code>lookupGrade()</code>	Matches a percentage to a <code>GradingScaleBand</code> , lower-inclusive/upper-exclusive (<code>minPercent <= p < maxPercent</code>) — a shared boundary (e.g. 80%) is never ambiguous; a percentage of exactly 100 matches the band whose <code>maxPercent</code> is 100. This boundary rule was corrected during verification after an inclusive-both-ends bug let a boundary score match two bands.
<code>computeUnweightedGPA()</code>	Plain arithmetic mean across whichever subjects have a resolvable <code>gradePoint</code> — no subject-credit/weighting concept anywhere in this schema. <code>null</code> (not <code>0</code>) if no subject has one.
<code>computeUnweightedAveragePercentage()</code>	The fallback ranking basis when no subject resolves a GPA (e.g. a label-only grading scale with no <code>gradePoint</code> anywhere) — the same unweighted-average shape, averaging each subject's percentage instead. Added specifically for Class Overview ranking.

GPA/percentage are never forced where the data doesn't support them: a marks-only framework produces a percentage but no grade/point; a descriptive-only framework produces neither — both fall out naturally from the same functions, with no special-casing.

(Source: [ASSESSMENT_RESULTS.md](#))

21. Report Card Architecture

`buildReportCard(studentId, audience)` assembles student/school/session/grade info, `fetchAssessmentResults()` 's subject results and GPA, and — reused directly, **not re-queried** — `fetchAcademicProgress()` 's attendance and evaluations. Rendered live at `/dashboard/report-card/[studentId]`.

Deliberately not modeled like `Certificate`. A certificate is a permanent, frozen-at-issuance snapshot; a Report Card must reflect corrections made after publication, so freezing one into a stored row would directly contradict the audited-correction design. **There is no persisted `ReportCard` model and no PDF export** — both explicitly out of scope.

Access: the Student themselves, a linked Parent, or staff (School Admin / any approved Teacher at the school, no assignment-level scoping — the same Skills-page precedent) — resolved to the matching audience before calling `buildReportCard()`, so a Student/Parent only ever sees published data through the identical filter used everywhere else.

(Source: [ASSESSMENT_RESULTS.md](#))

22. Promotion & Grade Decisions

`/dashboard/grades/[schoolGradeId]` — a School Admin opens one grade's Class Overview, multi-selects students from the roster, and applies one decision to the whole selected batch, entirely inside one `prisma.$transaction` (verified with a real 100-student batch: 328ms, `decided: 100, skipped: 0`).

Decision	status	outcomeGradeId
Promote	COMPLETED	required — the target grade
Repeat	REPEATED	required — same grade
Transfer	TRANSFERRED	none
Leave	LEFT	none

Eligibility is re-checked server-side: submitting an already-decided or foreign `gradeHistoryId` is silently excluded and counted in `skipped`, never double-decided. **Promotion is completely independent of section** — it never reads or writes `sectionId`; the roster's "Apply Decision" and "Assign Section" panels hit entirely separate endpoints.

Class Overview vs. Promotion-panel eligibility

The Class Overview roster uses the broadened `CURRENT_ROSTER_STATUSES` (see §12) so the page honestly shows everyone currently in the grade, decided or not. **The Promotion-action panel's own eligibility stays deliberately narrower** — only a still-`ENROLLED` row can be selected for a *new* decision, enforced independently server-side by the `grade-decisions` route regardless of what the roster displays. Broadening the display never changes who a decision can actually be applied to.

Class Overview additions

- **Teachers & Subjects** — every `TeacherAcademicAssignment` for the grade/session.
- **Section-wise grouping** — by the student's current-session `sectionId`; a final "Unassigned / No Section" group, never hidden; a per-section, **display-only** Roll No. (not a persisted field — see §37).
- **"Repeated" badge** — derived from the student's **prior-session** row only (`status === "REPEATED"` and its `outcomeGradeId` equals this grade), never from their current-session status — a forward-looking decision about *next* session is never conflated with how they arrived in *this* one.
- **Top 5 ranking** — computed **once, across the whole grade**, before any section grouping; reuses `fetchAssessmentResults()` / `computeUnweightedGPA()` / `computeUnweightedAveragePercentage()`, filtered to published results only.

Undecided students & the Pending/Unresolved queue

A student never decided stays **ENROLLED** indefinitely. On rollover, they're excluded from automatic placement and surface in a **persistent** Pending/Unresolved queue on [/dashboard/grades](#) — resolvable by recording the missing decision (audited) or manually placing them directly (unaudited, an honest gap in the record). Verified correct across a 3-session chain, including a student pending through an intervening session with zero rows of their own.

(Source: [GRADES_AND_PROMOTION.md](#))

23. Certificate System

`Certificate` keeps **recipient, instructor, and issuer as separate concepts**, never conflated — `recipientUserId`, optional `instructorId` (an `Instructor` can be named with no MEGA ID at all), `issuerType` (`MEGA_EDU` | `ORGANIZATION` | `SCHOOL` | `JOINT` — only `ORGANIZATION` is reachable today), and a separate `associatedSchoolId` for informational school context.

Issuance: `issueCourseCertificate()` (`src/lib/certificates.ts`) is the **only** code path that creates a `Certificate` for a course — called atomically alongside marking a `CourseEnrollment` complete, so an enrollment can never end up "complete" with no certificate or vice versa.

Grade certificates 🟡 — `Certificate.gradeHistoryId` is a reserved, unlinked column; `issueGradeCertificate()` **does not exist**. This was a deliberate Phase 2 scope exclusion, not an oversight — `GradeHistory` fully exists, but the issuance path was never built.

Display: `CertificateDocument.tsx`, a true-to-size A4 landscape layout driven by a pure view-model builder (`buildCertificateViewModel()`), reading only frozen snapshot fields plus a live logo lookup (the one deliberate exception to the snapshot rule — see §27). Shown at `/dashboard/certificates/[id]/preview` (access-gated to the recipient or a Platform Admin) and publicly, unstyled, at `/verify/[code]`.

Not built: PDF export 📄, QR code generation 📄 (a space is marked but empty). A school with a logo uploaded (see §24) shows it here; organizations still have no `logoUrl` field at all (see §37), so an organization-issued certificate always renders the name-only fallback — explicitly designed for, not a broken state.

(Source: [CERTIFICATES.md](#))

24. School Logos & User Profile Photos

An identity/media-upload layer added on top of the existing `School` and `User` models — no new models, one additive field on each.

Data model

- `School.logoUrl` (pre-existing field, previously unpopulated) — now actively managed. Displayed on School Directory cards, a school's public profile page, the homepage's "Recently joined schools," and the School Admin dashboard's identity header.
- `User.avatarUrl` (new field) — deliberately placed on the shared `User` model, **not** duplicated onto `Teacher`, `Student`, or `Parent`. This follows MEGA ID's core principle directly: **one identity, many roles** — a person's photo is a property of *them*, not of any one role they hold, so the same avatar automatically appears across every dashboard/list surface regardless of which role is being viewed.

API routes & authorization

Route	Methods	Authorization
<code>/api/schools/[id]/logo</code>	POST (upload/replace), DELETE (remove)	<code>requireSchoolAdmin(schoolId)</code> — the same helper every other school-profile write route uses
<code>/api/user/avatar</code>	GET (read own), POST (upload/replace), DELETE (remove)	Session-based self-service only — the caller's own <code>userId</code> , inline-checked the same way <code>/api/interests</code> already does. There is no route or code path through which one user can set another user's <code>avatarUrl</code> .

Upload validation

Enforced server-side in `src/lib/uploads.ts` — the client-supplied filename and browser-reported MIME type are never trusted:

- Accepted formats: **PNG, JPEG, WebP** — detected by reading the file's actual magic bytes, not the filename extension or the `Content-Type` header. **SVG is not accepted** (avoids script-injection risk from uploaded SVG markup).
- Maximum size: **2MB**.
- Stored filenames are always a fresh server-generated UUID (`crypto.randomUUID()`) plus the sniffed extension — the original uploaded filename is never used or persisted anywhere.
- A file that fails validation (wrong format, too large, or spoofed type) is rejected before anything is written to disk — no partial or orphaned file results from a rejected upload.

Storage architecture

Local filesystem, under `public/uploads/{schools,users}/...` (gitignored — never committed; user-uploaded content is never part of the repository). `logoUrl` / `avatarUrl` store the resulting root-relative URL, served by Next.js's normal static file handling and rendered through `next/image` for optimized, responsive delivery.

⚠ This storage implementation requires persistent filesystem storage and is appropriate for the current development/traditional server deployment model. Before deployment to typical serverless infrastructure, uploads should be migrated behind an object-storage adapter (e.g. S3-compatible) — a serverless/edge host's filesystem is typically ephemeral or read-only. No schema change would be needed for that migration; only the save/delete functions in `src/lib/uploads.ts` would change. See §31.

Replacement & removal lifecycle

Both upload routes follow the same strict ordering, so a failure partway through never leaves a user or school without their previous image:

1. The new image is validated and saved to disk first.
2. The database row (`logoUrl` / `avatarUrl`) is updated.
3. Only once both of those succeed is the *previous* stored file deleted.

A `DELETE` call removes the stored file and clears the field to `null`. Neither route accumulates unused files across normal replace/remove cycles — verified directly (upload → replace → confirm exactly one file remains and the old one is gone; remove → confirm the file is deleted from disk).

Fallback system — `src/components/Avatar.tsx`

One shared component renders both people and schools, with a deliberate visual distinction between the two rather than one generic treatment:

- **People** — a circular avatar. With no photo, a deterministic two-letter initials badge (first + last name initial) in one of the five brand accent colors, chosen by hashing the person's name — the same person always gets the same color, no randomness.
- **Schools** — a softly rounded-square "institution mark," bordered, always a fixed white-background/navy-text treatment (not the person palette) — both to read as distinctly institutional rather than personal, and so the fallback stays legible even placed on a navy background (e.g. the school profile page's hero band), where a randomly-assigned navy fallback would otherwise be invisible.

`/dashboard/profile` — the My Profile page

A new, deliberately minimal self-service page — available to any authenticated role — showing the caller's photo (with upload/replace/remove via `ProfilePhotoManager.tsx`), name, email, MEGA ID, linked school (where applicable), and role(s). Reachable from a small avatar/"My Profile" link in `SiteHeader.tsx`. Not a full account-management surface — no password change or editable name/email exist here; this is a foundation for future MEGA ID profile work, not that work itself.

(Source: `src/lib/uploads.ts`, `src/components/Avatar.tsx`, `src/app/api/schools/[id]/logo/route.ts`, `src/app/api/user/avatar/route.ts`, `src/app/dashboard/profile/page.tsx`)

25. Dashboard Architecture

`dashboard/page.tsx` is a single server component that branches on the caller's role-priority order (see §9), resolves the relevant data with direct Prisma queries, and renders the matching role-specific client component (`DashboardClient.tsx` , `TeacherDashboard.tsx` , `StudentDashboard.tsx` , `ParentDashboard.tsx` , `OrgDashboard.tsx` , `PlatformAdminDashboard.tsx`). Each role dashboard's `DashboardHero` now optionally shows an identity badge — the school's logo for the School Admin dashboard, the person's own avatar for Teacher/Student/Parent — via the shared `Avatar` component (see §24). A separate, role-agnostic `/dashboard/profile` page (also §24) is the self-service surface for managing that avatar, reachable from `SiteHeader.tsx` .

Shared presentational components, reused rather than duplicated per role:

- `AcademicProgressPanel.tsx` — Teaching Progress, Test Results, Recent Attendance, Teacher Evaluations, Assessment Results — parameterized by an explicit `audience: "STUDENT" | "PARENT" | "STAFF"` that filters what's shown (e.g. published-only for Student/Parent, unfiltered for Staff). Rendered once per Student, once per linked child on Parent, and on the Student Profile page.
- `MeetingActions.tsx` — all meeting create/complete/cancel/reschedule/link logic, reused identically on three surfaces (General Evaluations, Subject Evaluations panel, Meetings management page).

`fetchAcademicProgress()` / `fetchAssessmentResults()` / `fetchMeetingsForStudent()` (`src/lib/academicProgress.ts` , `src/lib/assessmentResults.ts`) are the shared query functions every dashboard branch and the Student Profile page call — never a second, parallel query implementation per surface.

(Source: [ARCHITECTURE.md](#), [ASSESSMENT_AND_EVALUATION.md](#), [ASSESSMENT_RESULTS.md](#))

26. API Architecture & Major Routes

All routes live under `src/app/api/**/route.ts`, grouped by resource. There is no separate backend — API routes and pages share the same Prisma client and `src/lib` helpers. Response bodies are JSON; a successful response generally includes `{ ok: true, ... }`, an error `{ error: string }`. **No GET list routes exist for any Phase 3 config/results model** — every read happens through the relevant page's own direct Prisma query, a deliberate, consistent convention across every admin/results page in this codebase.

The full, exact inventory (method, path, auth, request/response shape) is maintained in [API.md](#) and is not duplicated here to avoid drift between two copies of the same list. Summary by domain:

Domain	Route prefix	Auth
Auth & registration	<code>/api/auth/*</code>	none / NextAuth
Post-registration affiliation	<code>/api/teacher student parent/*</code> , <code>/api/schools/create-for-admin</code>	session
Platform Admin	<code>/api/admin/*</code>	<code>requirePlatformAdmin</code>
School directory & profile	<code>/api/schools/[id]</code> , <code>.../programs</code> , <code>.../news</code> , <code>.../opportunities</code> , <code>.../logo</code>	<code>requireSchoolAdmin</code>
Staff & students	<code>.../students</code> , <code>.../teachers</code> , <code>.../accountants</code>	<code>requireSchoolAdmin</code>
Sessions & grades (Phase 2)	<code>.../academic-sessions*</code> , <code>.../grades*</code> , <code>.../grade-placements</code> , <code>.../grade-decisions</code> , <code>.../section-assignments</code>	<code>requireSchoolAdmin</code>
Subjects & academic assignment (3A)	<code>.../subjects*</code> , <code>.../teacher-academic-assignments*</code>	<code>requireSchoolAdmin</code>
Operations (3B)	<code>.../class-teacher-assignments*</code> , <code>.../attendance*</code> , <code>.../units*</code> , <code>.../tests*</code>	<code>requireSchoolAdmin</code> OR <code>requireClassTeacher</code> / <code>requireTeacherAssignment</code>

Evaluation & PTM (3C)	<code>.../evaluations*</code> , <code>.../meetings*</code>	<code>requireSchoolAdmin</code> OR <code>requireClassTeacher</code> / <code>requireTeacherAssignment</code>
Assessment Framework (3D-1)	<code>.../grading-scales*</code> , <code>.../assessment-frameworks*</code> , <code>.../assessment-framework-assignments*</code>	<code>requireSchoolAdmin</code> only
Assessment Results (3D-2/3/4)	<code>.../components/[id]/results</code> , <code>.../publish</code> , <code>/api/schools/[id]/assessment-results/[resultId]</code>	<code>requireSchoolAdmin</code> OR <code>requireTeacherAssignment</code>
Organizations & courses	<code>/api/organizations/*</code> , <code>/api/courses/*</code> , <code>/api/enrollments/*</code>	<code>requireOrgAdmin</code> / <code>requireCourseOwner</code> /inline
Identity & notifications	<code>/api/interests/*</code> , <code>/api/notifications/*</code> , <code>/api/user/avatar</code>	session/inline

(Source: [API.md](#))

27. Business Rules & Data Integrity Protections

The full, individually-approved rule set — each with its rule, rationale, and applicability — is maintained in [PRODUCT_RULES.md](#). It is the tie-breaker whenever a future change seems to conflict with existing behavior: if a rule is documented there, it was a deliberate choice, not an oversight. Categories, summarized:

- **Data integrity & migration discipline** — no Prisma enums; additive-first migrations, verified before cleanup; the SQLite-specific bulk-write transaction pattern (flagged for Postgres rework); legacy fields retired by disuse, never deletion.
- **Snapshot fields vs. live lookups** — freeze display text on anything permanent (`Certificate` , `GradeHistoryAudit`); logos are the deliberate live-lookup exception; never hand-type or guess an asset — degrade gracefully to name-only.
- **The "one audited/gatekept write path" pattern** — `issueCourseCertificate()` , `recordGradeDecision()` , `reassignSection()` , `carryForwardEligibleStudents()` , `correctAttendance()` , `updateEvaluationRemarks()` , `correctComponentResult()` — each the sole writer for its concern.
- **New placements are creation, not decisions** — Initial Setup, rollover carry-forward, Add Student's inline placement, and the pending queue's manual placement are all unaudited row *creation*; only a change to an *existing* row is a decision requiring an audit.
- **Never guess when confidence is low** — `matchLegacyGradeText()` returns `null` rather than a low-confidence guess; undecided students are never silently defaulted at rollover.
- **The recurring `NULL ≠ NULL` gap** — documented as a standing reminder: any new model with an optional field in its unique key needs an explicit application-level pre-check, not just the database constraint.
- **Per-phase rules** for Subjects/Teacher Assignment (3A), Academic Operations (3B), Evaluation/PTM (3C), Assessment Framework (3D-1), and Assessment Results/Publishing (3D-2/3/4) — each documented individually in [PRODUCT_RULES.md](#) with rule, rationale, and verification evidence.
- `CURRENT_ROSTER_STATUSES` and the **Promotion-panel's deliberately narrower eligibility** — the newest additions, documented in full in [PRODUCT_RULES.md](#) and summarized in §12/§22 above.

28. Audit Trails

Model audited	Audit table	Written by	Trigger
<code>GradeHistory</code> (status/outcome)	<code>GradeHistoryAudit</code>	<code>recordGradeDecision()</code>	Every decision, including the first ever made on a row
<code>GradeHistory</code> (section)	<code>GradeHistoryAudit</code> (same table, <code>previousSectionId</code> / <code>newSectionId</code>)	<code>reassignSection()</code>	Every reassignment on an <i>existing</i> row (not the initial set-at-creation value)
<code>Attendance</code>	<code>AttendanceAudit</code>	<code>correctAttendance()</code>	Every correction, capturing both status and remarks even if only one changed
<code>StudentEvaluation</code>	<code>StudentEvaluationAudit</code>	<code>updateEvaluationRemarks()</code>	Only once shared with either Parent or Student — silent while still a private draft
<code>AssessmentComponentResult</code>	<code>AssessmentComponentResultAudit</code>	<code>correctComponentResult()</code>	Only once the parent subject's publication is <code>PUBLISHED</code> — silent while still <code>DRAFT</code>

Deliberately *not* audited (current-state, operational, freely re-creatable data, not permanent decisions): `TeacherGradeAssignment` , `GradeSubject` , `TeacherAcademicAssignment` , `ClassTeacherAssignment` , `AssessmentFramework` / `AssessmentPeriod` / `AssessmentComponent` / `GradingScale` / `AssessmentFrameworkAssignment` (structural config, not history), `ParentTeacherMeeting.outcomeNotes` and rescheduling.

Every audit table in this system is genuinely **append-only** — no route anywhere updates or deletes an audit row.

29. Security Model

- **Session:** JWT-based (NextAuth), not database sessions; `bcrypt.compare` for login.
- **Authorization:** enforced server-side on every write route via the `requireX()` suite (see §8) — never trusted from UI state alone. Verified repeatedly, throughout every phase, via direct `fetch()` calls from an unauthorized session confirming `403`, not just that a button was hidden.
- **Data isolation:** every `studentId` / `schoolId` used in a read is resolved server-side from the caller's own session-derived relationships (a Parent's own `parent.children`, a School Admin's own `SchoolAdmin` row) — never trusted from client-supplied query parameters. Verified live for Parent academic visibility, the Class Overview roster, and the Student Profile page.
- **What's absent** 🚫: no rate limiting on login/registration/any route; no CSRF protection beyond NextAuth's own; no session revocation ("log out everywhere" — JWTs live until expiry); no OAuth/SSO; no email verification; no password reset flow; no audit log of authorization *decisions* themselves (as distinct from the data-level audit trails in §28).

(Source: [AUTHENTICATION_AND_AUTHORIZATION.md](#), [MEGA_ID.md](#), [KNOWN_GAPS.md](#))

30. Calculation Engines & Shared Libraries

Module	Purpose
<code>auth.ts</code>	NextAuth configuration
<code>authorize.ts</code>	The <code>requireX()</code> guard suite — see 58
<code>prisma.ts</code>	Singleton Prisma client
<code>notify.ts</code>	<code>notify()</code> / <code>notifySchoolCommunity()</code> — best-effort, never allowed to fail the calling action
<code>gradeHistory.ts</code>	<code>recordGradeDecision()</code> , <code>reassignSection()</code> , <code>CURRENT_ROSTER_STATUSES</code>
<code>gradeMatching.ts</code>	<code>matchLegacyGradeText()</code> — never guesses
<code>gradeRollover.ts</code>	<code>carryForwardEligibleStudents()</code> , <code>findPendingStudents()</code> — idempotent, re-runnable
<code>attendance.ts</code>	<code>correctAttendance()</code>
<code>evaluation.ts</code>	<code>updateEvaluationRemarks()</code> , <code>shareEvaluation()</code>
<code>assessmentFramework.ts</code>	<code>resolveFrameworkAssignment()</code> , <code>assignmentCollisionExists()</code> , <code>componentCollisionExists()</code>
<code>assessmentResults.ts</code>	<code>computeComponentContribution()</code> , <code>aggregateGroup()</code> , <code>lookupGrade()</code> , <code>computeUnweightedGPA()</code> , <code>computeUnweightedAveragePercentage()</code> , <code>fetchAssessmentResults()</code> , <code>correctComponentResult()</code> , <code>buildReportCard()</code>
<code>academicProgress.ts</code>	<code>fetchAcademicProgress()</code> , <code>fetchMeetingsForStudent()</code> — the shared, audience-parameterized query functions behind every dashboard branch and the Student Profile page
<code>certificates.ts</code> / <code>certificateView.ts</code>	<code>issueCourseCertificate()</code> , <code>buildCertificateViewModel()</code>

All five "sole write path" functions (`issueCourseCertificate` , `recordGradeDecision` , `reassignSection` , `carryForwardEligibleStudents` , `correctAttendance`) share the same shape: typed input, optional `tx?: Prisma.TransactionClient` for composing into a larger transaction, transactional by default otherwise.

(Source: [ARCHITECTURE.md](#))

31. Deployment & Setup Requirements

Local development

```
npm install
npm run dev      # next dev - http://localhost:3000
```

Other scripts: `npm run build`, `npm run start`, `npm run lint`, `npm run db:push` (`prisma db push` — direct schema apply, no migration files), `npm run db:seed` (`tsx prisma/seed.ts`, idempotent), `npm run db:studio`.

Database

Development uses SQLite (`prisma/dev.db`). PostgreSQL is the marked, **never-configured-or-tested** production target. Switching requires changing `provider = "sqlite"` to `provider = "postgresql"` in `schema.prisma` and pointing `DATABASE_URL` at a real connection string.

File uploads — School Logos & Profile Photos ⚠️

Saved to the local filesystem under `public/uploads/` (see §24). **This storage implementation requires persistent filesystem storage and is appropriate for the current development/traditional server deployment model. Before deployment to typical serverless infrastructure, uploads should be migrated behind an object-storage adapter.**

Production 🔭

Nothing has been deployed. No hosting platform, no CI/CD pipeline, no Dockerfile, no production `next.config.js` overrides exist in this repository — genuinely unstarted work, not an oversight in documentation.

Known deployment requirements, inferred from the codebase (not yet acted on)

1. A Postgres database, with the schema `provider` switched and `npx prisma db push` (or a proper migration) run against it.
2. A real `NEXTAUTH_SECRET` / `NEXTAUTH_URL` matching the deployed domain.
3. A real `SEED_ADMIN_EMAIL` / `SEED_ADMIN_PASSWORD` before seeding, or skip seeding in production entirely.
4. **Review the SQLite-specific bulk-write transaction pattern before relying on it under Postgres** (see below) — this is the single most concrete pre-deployment code risk currently documented.

5. Standard Node.js runtime requirements for a Next.js 14 App Router app — no edge-specific code exists anywhere in this codebase.

(Source: [DEPLOYMENT.md](#))

32. Environment Variables

The complete, real list — nothing else is read anywhere in the app:

Variable	Purpose	Dev default
<code>DATABASE_URL</code>	Prisma connection string	<code>file:./dev.db</code>
<code>NEXTAUTH_SECRET</code>	JWT signing secret	placeholder — must be replaced for any real deployment
<code>NEXTAUTH_URL</code>	Canonical app URL for NextAuth callbacks	<code>http://localhost:3000</code>
<code>SEED_ADMIN_EMAIL</code>	Platform Admin account created by <code>db:seed</code>	<code>admin@megaedu.local</code>
<code>SEED_ADMIN_PASSWORD</code>	Same	<code>ChangeMe123!</code> — must be changed before seeding a real environment

(Source: [DEPLOYMENT.md](#))

33. Database Migration & Update Procedures

This project uses `npx prisma db push` — direct schema application, **no** `prisma migrate / migrations/ folder` — appropriate for the current single-environment SQLite setup, not necessarily for a multi-environment production deployment.

Established discipline (both Phase 1 and Phase 2 followed the same sequence, enforced manually, not by tooling):

1. Add new models/fields **additively** — nothing existing removed in the same pass.
2. Before applying anything that could conflict with existing data (a new `@@unique` constraint, for example), write a one-off verification script, run it, confirm no conflicts, delete the script.
3. Apply with `npx prisma db push`.
4. Only after a change is applied and verified does any "clean up the old thing" pass happen — and some legacy fields (`Student.gradeLevel`) are intentionally **never** dropped.

⚠ The Postgres migration caveat, specifically

`grade-placements` and `teacher-assignments` catch a unique-constraint violation (`P2002`) *inside* an open transaction and continue looping — verified to work correctly on SQLite (a caught statement error doesn't poison the rest of the transaction), but this is **not true on PostgreSQL**, where a failed statement aborts the whole transaction until an explicit rollback. This must be re-checked (likely reworked to pre-filter duplicates before the transaction opens) before any Postgres migration. **Not applicable** to `grade-decisions` or the rollover sweep, which validate eligibility *before* opening their transaction — a genuinely different, already Postgres-safe pattern.

(Source: [ARCHITECTURE.md](#), [PRODUCT_RULES.md](#), [DEPLOYMENT.md](#))

34. Development Workflow & Verification

This project has **no automated test suite** (no Jest/Vitest/Playwright/Cypress, no `*.test.*` files anywhere). The established substitute, followed consistently across every phase of this project, is a disciplined manual workflow:

1. **Investigate before designing.** A read-only pass over the relevant schema, `authorize.ts` patterns, and any precedent feature (e.g. reading `certificateView.ts` before designing the Report Card, or `UnitTestResult`'s eager-roster pattern before deciding Assessment Results should be lazy instead) precedes any code.
2. **Design and get explicit approval before implementing.** Non-trivial features (a new phase, a schema change) get a proposed design reviewed and approved before any code is written — the same discipline this very documentation task followed (Table of Contents approved before writing).
3. **Implement additively**, following the migration discipline in §33.
4. **Typecheck as a gate:** `npx tsc --noEmit -p tsconfig.json` after every non-trivial change.
5. **Verify live, against the running dev server and the real database** — not mocked, not assumed. This includes negative cases (confirming an unrelated user is correctly redirected/403'd, not just that the right person succeeds) and, for anything transactional, confirming the transaction actually behaves atomically (a duplicate is skipped without crashing the batch; a partial failure doesn't leave inconsistent rows).
6. **Throwaway verification scripts**, run once via `npx tsx` and then deleted — never left behind in the repository. Used for anything touching real data: checking for existing duplicate rows before adding a constraint, timing a bulk route, proving a sweep is idempotent.
7. **Clean up test-data afterward, with before/after row-count confirmation.** Every phase's verification pass ends by confirming the school's real, pre-existing data (student/teacher/assignment counts) was unaffected by the throwaway fixtures used during testing.
8. **Update documentation in the same pass**, not as an afterthought — if a change alters an architectural or business-rule decision, the relevant `/docs` file is updated (or, for a significant behavior change, fully rewritten) in the same body of work, distinguishing  /  /  /  status accurately rather than blurring them to make a feature sound more finished than it is.
9. **Get explicit approval before committing/pushing.** Git history is not touched (`git add` / `git commit` , let alone `push`) until the user has reviewed the change and explicitly asked for it.

Seeded demo accounts (`prisma/seed.ts` , idempotent) are the de facto manual test fixtures — one account per role, plus a verified demo school/course. See [TESTING.md](#) for the full account list and the verification evidence recorded for every phase.

(Source: [DEVELOPMENT_GUIDELINES.md](#), [TESTING.md](#))

35. Demo & Sample Environment

Beyond the minimal bootstrap seed described above, `prisma/seed-demo.ts` builds an optional, idempotent, self-sufficient demonstration environment on top of it — for testing, demonstration, training, exploring workflows, and multi-school testing.

Two demonstration schools, deliberately different in scale, to demonstrate genuine multi-school data isolation:

- **Sunrise Academy** — the flagship, full-depth demonstration environment.
- **Himalayan Secondary School** — a smaller, independent school, showing that each school's data (sessions, grades, teachers, students) stays fully isolated from the other.

The environment includes realistic **fictional** examples of schools, teachers, students, parents, academic structures, promotion history, assessment results, attendance, evaluations, unit tests, Parent-Teacher Meetings, and courses/certificates — all generated deterministically (a seeded PRNG), so re-running the seed produces identical results every time.

```
npm run db:seed          # baseline bootstrap
npm run db:seed:demo     # full demo environment
npm run db:verify:demo  # confirms internal consistency against real calculation logic
```

Full operational detail — the complete account/login list, reset procedure, and what each module contains — is maintained in [DEMO_DATA.md](#) and not duplicated here.

(Source: [DEMO_DATA.md](#))

36. Maintenance Guidelines

For anyone — human or AI-assisted — making a future change to this codebase:

- **Read `PRODUCT_RULES.md` before touching business logic.** If a change seems to conflict with a documented rule, that's a signal to stop and check, not to route around it.
- **Read `ARCHITECTURE.md` before touching architecture.** Match the established conventions (server components fetch Prisma directly; the "one audited/gatekept write path" pattern; `router.refresh()` as the only revalidation strategy) rather than introducing a new pattern in one corner of the app.
- **Scope discipline:** a bug fix doesn't need surrounding cleanup; a one-shot operation doesn't need a speculative helper. If a request describes 5 steps, build 5 steps.
- **Data integrity:** preserve historical records — no route should delete `GradeHistory`, `GradeHistoryAudit`, or `Certificate` rows. Test migrations for data conflicts before applying a new constraint. Never use a Prisma `enum`.
- **Don't invent.** If a request assumes something exists that a codebase search doesn't confirm (the documented example: an earlier request assumed "premium courses"/"course bundles" were previously approved — a direct search found zero evidence they were ever built or designed), say so rather than building a plausible-sounding version to satisfy the request's framing.
- **Verify, don't assume.** Typecheck, then exercise the actual change against the running dev server and real database with throwaway fixtures, cleaned up afterward.
- **Update documentation in the same pass** as any architectural or business-rule change, keeping the  /  /  /  distinction honest.

(Source: [DEVELOPMENT_GUIDELINES.md](#))

37. Known Gaps & Deliberate Out-of-Scope Decisions

The complete, individually-re-verified list is maintained in [KNOWN_GAPS.md](#). Organized here by category and status:

Known limitations (real gaps, not by design)

Gap	Detail
<code>School.isActive</code> / <code>Organization.isActive</code> never set to <code>false</code>	Read in two places, no deactivation action exists anywhere
Organization verification not enforced	Nothing checks <code>Organization.verified</code> before course publishing or enrollment
SQLite-specific bulk-write transaction pattern	Would misbehave on PostgreSQL — see §33
No automated test suite	See §34 for the manual substitute
<code>Organization</code> has no <code>logoUrl</code> field	Certificates fall back to name-only for organization-issued certificates
"Roll No." on Class Overview is display-only	Not a persisted schema field — a per-section sequential position computed at render time
School logo / avatar uploads assume a persistent local filesystem	Appropriate for the current dev/traditional-server model; would need an object-storage adapter before serverless deployment — see §24, §31

Deliberately out of scope (approved decisions, not oversights)

Area	Deferred item
Sections	Section-level teacher assignment, section-level analytics/reporting
Subjects/Units	Copying a <code>GradeSubject</code> offering or <code>TeachingPlan</code> / <code>TeachingUnit</code> set forward from a prior session
Teaching structure	Any teaching hierarchy (primary/assistant/substitute teacher)
Unit Tests	Retest concept
Evaluations/PTM	Un-share path for evaluations; parent-initiated meeting requests; meeting recurrence/series

Assessment	Subject-credit/weighting concept (GPA stays unweighted); <code>UnitTest</code> marks auto-derivation; <code>GET</code> list API routes for any Phase 3D model
Promotion	<code>GradingScaleBand.isPassing</code> is not read by any promotion logic — reserved for a future reference display only
Certificates	Grade-certificate issuance (<code>issueGradeCertificate()</code> doesn't exist); PDF export; QR code generation
Commerce	Payment processor integration; premium/bundle course concepts (confirmed via direct search: zero evidence they were ever designed)
Auth	OAuth/SSO, email verification, password reset, rate limiting, session revocation, account deactivation/deletion
Platform Admin	No in-app route to grant/revoke <code>PLATFORM_ADMIN</code>

● Designed but not implemented

- Grade-based certificates (a reserved field exists on `Certificate` , but no issuance path).

(Source: [KNOWN_GAPS.md](#))

38. Appendix: Documentation Index

Document	Covers
PROJECT_OVERVIEW.md	High-level status of every phase, module list, current dev status table
ARCHITECTURE.md	Rendering patterns, <code>src/lib</code> "only path" pattern, module relationships
DATABASE.md	Full field-by-field Prisma model inventory
API.md	Complete API route inventory with request/response shapes
AUTHENTICATION_AND_AUTHORIZATION.md	Login, session, the full <code>requireX()</code> suite
USER_ROLES.md	What every role can do, approval workflows
MEGA_ID.md	The identity model and its design principles
ACADEMIC_SESSIONS.md	Sessions, rollover, the one-ACTIVE-per-school rule
GRADES_AND_PROMOTION.md	Grades, sections, promotion, Class Overview
ACADEMIC_STRUCTURE.md	Subjects, <code>GradeSubject</code> , <code>TeacherAcademicAssignment</code> (Phase 3A)
ACADEMIC_OPERATIONS.md	Class/Section Teachers, Attendance, Teaching Units, Unit Tests (Phase 3B)
ASSESSMENT_AND_EVALUATION.md	Evaluations, Parent-Teacher Meetings (Phase 3C)
ASSESSMENT_FRAMEWORK.md	Framework/grading-scale configuration, the guided wizard (Phase 3D-1)
ASSESSMENT_RESULTS.md	Marks entry, publishing, calculation engine, Report Card (Phase 3D-2/3/4)
CERTIFICATES.md	Certificate model, issuance, display
COURSES_AND_ENROLLMENTS.md	MEGA Academy courses and enrollment
PRODUCT_RULES.md	Every explicitly-approved business rule, with rationale
KNOWN_GAPS.md	Every known gap and deliberate scope decision, individually re-verified
CHANGELOG.md	Dated, hand-written record of every notable change

DEPLOYMENT.md	Local dev setup, environment variables, production readiness
DEVELOPMENT_GUIDELINES.md	Rules for future work on this codebase
TESTING.md	Verification practice, seeded demo accounts
DEMO_DATA.md	The full demo/sample environment: schools, accounts, login info, seed/reset/verify commands

Note on the prior technical PDF: [docs/MEGA_EDU_Technical_Documentation.pdf](#) predates Phases 3A–3D and the guided wizard/Class Overview work (last touched 2026-08-28). This Markdown document supersedes it as the current technical reference; the PDF has been left in place rather than deleted, since removing it was outside the scope of this task.