

MEGA.EDU — Technical Documentation

Audience: developers, technical team members, system administrators, and future maintainers. **Status legend** (used throughout): ■ **Implemented** · ■ **Designed/approved, not yet implemented** · ■■ **Known gap/issue** · ■ **Future/planned** **Last verified:** 2026-09-10 (K1-K8 reconciliation — Organization provider profile, Academy participation/global visibility, user-centric enrollment, Opportunity edit/delete, course↔provider links, MEGA Academy navigation label, Organization logo, Organization Events & Resources, obsolete certificate backfill script removal — added on top of Calendar Kilometer 1/1.1/1.2, My Profile Kilometer 1, Student Profile Academic Snapshot), against the current codebase and the audited /docs documentation set. **Source discipline:** every claim in this document is drawn from the existing, individually-audited documents in /docs (cross-referenced throughout) and, where a doc was ambiguous, from direct inspection of the implementation. Nothing here describes a planned or hypothetical feature as if it were built. Where something is designed but not implemented, or implemented but deliberately incomplete, that is stated explicitly.

Table of Contents

1. [Introduction & Purpose of This Document](#)
2. [System Overview](#)
3. [System Data Flow](#)
4. [Module Dependency Map](#)
5. [Technology Stack](#)
6. [Project Structure](#)
7. [Database Architecture & Prisma Models](#)
8. [Authentication & Authorization](#)
9. [User Roles & Permissions](#)
 - [School & Multi-School Architecture](#)
 - [10a. Institutional Affiliation & Context Architecture](#)
11. [Academic Sessions, Grades, Sections & Subjects](#)
12. [Student Lifecycle & GradeHistory](#)
13. [Teacher Assignments & Authorization Scopes](#)
14. [Attendance System](#)
15. [Evaluations](#)
16. [Parent–Teacher Meetings](#)
 - [Unit Tests](#)
 - [17a. Homework — Phase 1](#)
 - [17b. Calendar — Kilometer 1 / 1.1 / 1.2](#)
18. [Assessment Framework System](#)
19. [Assessment Results & Publishing Workflow](#)
20. [Grading & GPA Calculations](#)
21. [Report Card Architecture](#)
22. [Promotion & Grade Decisions](#)
23. [Certificate System](#)
24. [School Logos & User Profile Photos](#)
25. [Dashboard Architecture](#)
26. [API Architecture & Major Routes](#)

27. [Business Rules & Data Integrity Protections](#)
 28. [Audit Trails](#)
 29. [Security Model](#)
 30. [Calculation Engines & Shared Libraries](#)
 31. [Deployment & Setup Requirements](#)
 32. [Environment Variables](#)
 33. [Database Migration & Update Procedures](#)
 34. [Development Workflow & Verification](#)
 35. [Demo & Sample Environment](#)
 36. [Maintenance Guidelines](#)
 37. [Known Gaps & Deliberate Out-of-Scope Decisions](#)
 38. [Appendix: Documentation Index](#)
-

1. Introduction & Purpose of This Document

This document is a single technical reference for MEGA.EDU, consolidating the detailed, individually-verified documents already maintained under `/docs` (each covering one subsystem in depth) into one narrative a new developer, system administrator, or maintainer can read top to bottom. It does not replace those documents — it cross-references them throughout, and the reader is expected to follow those links for full implementation detail (exact request/response shapes, full verification evidence, field-by-field schema notes). This document is the map; the individual `/docs` files are the territory.

Everything below reflects the system **as implemented and verified**, as of the phases completed through **Phase 3D-2/3/4** (Assessment Results, Publishing, Report Cards) plus the subsequent **guided Assessment Wizard**, **Class Overview** enhancements, the **School Logos & User Profile Photos** identity system, **Homework**, **Phase 1** ([§17a](#)), and, most recently, **Calendar**, **Kilometer 1/1.1/1.2** ([§17b](#)) and **My Profile**, **Kilometer 1** ([§24](#)). Nothing here is invented, and nothing planned is described as available today — see [§37](#) for what is deliberately out of scope versus genuinely missing.

2. System Overview

MEGA.EDU is a national education network connecting schools, teachers, students, parents, and education organizations under one identity system, **MEGA ID**. A school gets a public profile and a digital identity; teachers and students at that school get accounts a school administrator approves; organizations can publish free online courses (MEGA Academy) and post opportunities; parents can follow their children's academic progress. Everyone signs in with the same account type, regardless of role.

The platform's purpose is to give every school, teacher, and student a **verifiable digital presence**: a school's identity isn't tied to one administrator's login, a student's academic record and certificates follow *them* rather than staying locked inside one school's paperwork, and achievements are independently verifiable by anyone with the right link.

High-level architecture

MEGA.EDU is a **single Next.js 14 App Router application** — there is no separate backend service. Pages and API routes are colocated under `src/app`, both talking to the same Prisma client against one database.

(Diagram — see the corresponding .svg file under docs/assets/, or the Markdown source's Mermaid block, for the full diagram; omitted from this PDF export.)

Technical Mermaid Source (for future editing) *(Diagram source omitted from PDF export — see the Markdown source or the rendered .svg asset.)*

(Source: [ARCHITECTURE.md](#))

There is no client-side global state manager. `router.refresh()` — re-running the server component and re-fetching its data — is the application's entire "revalidate after mutation" strategy, following a `fetch()` → `check res.ok` → `router.refresh()` mutation flow throughout.

3. System Data Flow

The diagram below traces the major academic and assessment data flows a real school follows during a session — from setup through to a published Report Card and a promotion decision. It complements, rather than duplicates, the [Module Dependency Map](#) below: this shows the *order things happen in*, the map shows *what depends on what structurally*.

(Diagram — see the corresponding .svg file under docs/assets/, or the Markdown source's Mermaid block, for the full diagram; omitted from this PDF export.)

Technical Mermaid Source (for future editing) *(Diagram source omitted from PDF export — see the Markdown source or the rendered .svg asset.)*

Reading this diagram: setup (session, grades, sections, subjects, teachers, students) must happen before any daily operations can run, since attendance, evaluations, units, and assessment entry are all scoped to a `GradeHistory/TeacherAcademicAssignment` combination that must already exist. Assessment results and Report Cards are the only flow that terminates in something a Parent/Student ever sees directly — every other flow (attendance, evaluations, unit tests) is visible to Student/Parent as read-only progress data, not a publish-gated artifact. Promotion is the flow that closes the loop back to a new `GradeHistory` row for the next session.

4. Module Dependency Map

This map shows structural dependencies between the major Phase 3 modules — an arrow means "the module at the arrow's head cannot function without data or authorization state from the module at its tail."

(Diagram — see the corresponding .svg file under docs/assets/, or the Markdown source's Mermaid block, for the full diagram; omitted from this PDF export.)

Technical Mermaid Source (for future editing) *(Diagram source omitted from PDF export — see the Markdown source or the rendered .svg asset.)*

Solid arrows are hard, load-bearing dependencies (the dependent module's routes or queries directly reference the depended-on module's rows). Dotted arrows are looser, informational relationships (e.g. a `ParentTeacherMeeting` may link to an `StudentEvaluation`, but neither requires the other to exist).

Key structural facts this map summarizes:

- **GradeHistory (Promotion/PROMO) underlies everything.** Attendance, Unit Tests, Evaluations, and Assessment Results all resolve a student's current placement through `GradeHistory`, not an independently-maintained roster.
- **TeacherAcademicAssignment/ClassTeacherAssignment gate almost every teacher-facing write.** Attendance and general Evaluations require `requireClassTeacher()`; Unit Tests, Subject Evaluations, and Assessment Results marks entry all require `requireTeacherAssignment()` scoped to the exact subject.
- **Assessment Results is the only module both Report Card and Class Overview ranking depend on.** Both reuse `fetchAssessmentResults()/buildReportCard()` — see [§20](#) and [§30](#).
- **Promotion (recordGradeDecision()) has no dependency on Assessment Results at all** — a Promotion decision can be recorded with zero assessment data ever entered. The relationship only exists in the *display* direction: Class Overview reads both the roster (from `GradeHistory`) and the ranking (from Assessment Results) on the same page, but neither module's write path touches the other.

5. Technology Stack

Layer	Technology	Notes
Framework	Next.js 14 (App Router)	Single deployment, pages and API routes colocated
UI	React 18, TypeScript	No separate frontend framework or SPA build
Styling	Tailwind CSS	Utility-first, no CSS-in-JS library
ORM	Prisma 5.20	Single <code>PrismaClient</code> singleton (<code>src/lib/prisma.ts</code>)
Database (dev)	SQLite	Single file, <code>prisma/dev.db</code>
Database (production, planned)	PostgreSQL ■	Never configured or tested — see §31
Authentication	NextAuth 4	Single Credentials provider, JWT session strategy
Password hashing	<code>bcryptjs</code>	<code>bcrypt.compare/bcrypt.hash(password, 10)</code>
Validation	<code>zod</code>	Used on registration routes
Dev tooling	TypeScript compiler, ESLint (<code>next lint</code>), Prisma Studio	No automated test framework installed — see §34

(Source: [PROJECT_OVERVIEW.md](#), [DEPLOYMENT.md](#))

No Prisma enums are used anywhere in this schema — SQLite's Prisma connector doesn't support them, even unused ones. Every status/type/role field is a plain `String`, with valid values documented in a comment above the field. This is a project-wide, explicitly-approved rule (see [§27](#)), not an oversight, and it is intended to remain the convention even after a future move to PostgreSQL (which does support enums).

6. Project Structure

```
src/ app/ # pages (App Router) + API routes, colocated api/ # route.ts handlers, grouped by
resource dashboard/ # role-specific dashboard pages/components setup/ # Initial School
Setup wizard (Phase 2) grades/ # Class Overview / Promotion rosters + Pending queue
sessions/new/ # New Session rollover academics/ # Subject catalog, grade offerings, teacher
assignments, # [gradeSubjectId]/ – Teaching Plans, Units, Tests attendance/ # Daily
attendance marking + correction evaluations/ # General Student Evaluations + meeting
scheduling meetings/ # Cross-role Parent-Teacher Meetings management students/[studentId]/
# Staff-only Student Profile page profile/ # Self-service "My Profile" – avatar, identity
info assessment-frameworks/ # Landing page, guided wizard, Advanced management
assessment-results/ # Marks entry, publish workflow report-card/[studentId]/ # Live Report
Card view admin/ # Platform Admin-only pages ... # public pages: schools, courses,
opportunities, etc. components/ # shared client components (SiteHeader, DashboardHero, #
AcademicProgressPanel, MeetingActions, certificate/) lib/ # server-side helpers – see §30
(includes uploads.ts) types/ # ambient type augmentation (next-auth.d.ts) prisma/
schema.prisma # single source of truth for the data model seed.ts # idempotent bootstrap
fixtures (upserts) seed-demo.ts # optional, fuller demo/sample environment – see §35
```

(Source: [ARCHITECTURE.md](#))

Rendering pattern

Two consistent patterns are used throughout, never mixed within one page:

1. **Server component fetches Prisma directly, renders (or passes props to) a client component.** The dominant pattern for a logged-in user's own dashboard data — `dashboard/page.tsx` branches on role, runs the relevant Prisma query, hands a fully-hydrated object to a `"use client"` component. Every Phase 2/3 configuration and roster page follows this.
2. **Client component fetches its own API route.** Used for mutations and independently-pollled data (e.g. the notification bell). The mutation flow is consistently `fetch()` → `check res.ok` → `router.refresh()`.

The `src/lib` layer — "the only path" pattern

A recurring, deliberate architectural decision: for any write that must be atomic with a side effect, or must never be bypassed, the codebase centralizes it into one function, and every caller goes through it rather than calling `prisma.<model>.update()` directly. See [§30](#) for the full inventory.

7. Database Architecture & Prisma Models

Datasource: SQLite in development (`prisma/dev.db`); the schema's own header comment and `.env.example` both mark PostgreSQL as the intended production target — nothing production-specific is configured yet. Every model below is ■ implemented (exists, migrated, and has at least one route reading/writing it) unless marked otherwise. Full field-by-field detail lives in [DATABASE.md](#); this section is an organized summary.

MEGA ID & roles

- **user** — the single identity record (email unique, `passwordHash`, `name`, `avatarUrl`). Every other model traces back to a `User` for who did what. `avatarUrl` lives here rather than on `Teacher/Student/Parent` specifically because a photo belongs to the person, not any one role — see [§24](#).

- **UserRole** — which role(s) a user holds (PLATFORM_ADMIN | SCHOOL_ADMIN | TEACHER | STUDENT | PARENT | ORGANIZATION_ADMIN | ACCOUNTANT), @@unique([userId, role]) — a user can hold several different roles at once.

Schools & Organizations

- **School** — verified/isActive flags (■■ isActive is read in two places but nothing ever sets it false — no deactivation action exists), logoUrl (uploadable/manageable by an authorized School Admin — see §24).
- **SchoolAdmin/SchoolAccountant** — join tables granting access to one school.
- **Organization, OrganizationAdmin/OrganizationAccountant** — the parallel structure for MEGA Academy publishers and Opportunity/Event/Resource providers. verified/isActive flags (same isActive-never-written caveat as School), independent academyParticipant (self-service MEGA Academy participation toggle — see §23), and logoUrl (K6 — uploadable/manageable by an Organization Admin, reusing the exact same upload infrastructure as School's logoUrl — see §24). No institutional-context history layer exists yet (OrganizationAdmin/OrganizationAccountant remain flat join tables — see §10a).

People

- **Teacher, Student** — optional 1:1 with User, optional schoolId, approved flag. Student.gradeLevel is the **permanent legacy free-text grade fallback** — retained forever, never dropped, no longer written once a school completes Initial Setup.
- **Parent, ParentStudent** — a parent's linked children, matched by the child's email.

Academic Sessions & Grades (Phase 2)

- **AcademicSession** — one school-year window; at most one ACTIVE per school (app-enforced).
- **GradeReference** — the platform-wide, fixed 13-row grade ladder (PP1–PP3, Y1–Y10); not school-editable.
- **SchoolGrade** — a school's opt-in to one GradeReference, with its own displayName.
- **Section** — an optional, school-defined subdivision of a SchoolGrade; not session-scoped; soft-deactivate only, no hard delete.
- **TeacherGradeAssignment** — per-session teacher-to-grade link.
- **GradeHistory** — a student's grade placement for one session, the permanent record; status: ENROLLED | COMPLETED | REPEATED | TRANSFERRED | LEFT. See §12.
- **GradeHistoryAudit** — append-only log of every decision and section change ever written to a GradeHistory row.

Subjects & Teacher Academic Assignment (Phase 3A)

- **Subject** — school-wide catalog, reusable across every grade/session; deactivate only.
- **GradeSubject** — which subjects a grade offers, **for one specific academic session** — never carried forward.
- **TeacherAcademicAssignment** — teacher → session → grade → optional section → subject; a real DELETE route exists (current-state, not historical).

School Academic Operations (Phase 3B)

- **ClassTeacherAssignment** — Grade Coordinator (sectionId: null) or Class Teacher; grade-wide and section-specific may coexist (unlike TeacherAcademicAssignment).
- **Attendance/AttendanceAudit** — one row per student per calendar day; audited corrections. See §14.
- **TeachingPlan/TeachingUnit** — a declared planned-total (separate model, not inferred from unit count) and the actual curriculum units with a teaching-progress status.
- **UnitTest/UnitTestResult** — pre-created roster pattern. See §17.

Teacher Qualitative Evaluation & Parent-Teacher Meetings (Phase 3C)

- `StudentEvaluation/StudentEvaluationAudit` — General vs. Subject via one nullable field; audit-on-share. See [§15](#).
- `ParentTeacherMeeting` — periodic and occasional through one model. See [§16](#).

Homework (Phase 1)

- `Homework` — one teacher-authored item per grade (or one section)/subject/session; no per-student fan-out, `teacherId` → `Teacher` (not `createdByUserId` → `User`). See [§17a](#).

Assessment Framework Foundation (Phase 3D-1)

- `AssessmentFramework`, `AssessmentPeriod`, `AssessmentComponent`, `GradingScale`, `GradingScaleBand`, `AssessmentFrameworkAssignment`. See [§18](#).

Assessment Results, Publishing, Report Cards (Phase 3D-2/3/4)

- `AssessmentComponentResult`, `AssessmentComponentResultAudit`, `AssessmentResultPublication`, plus `GradingScaleBand.isPassing`. See [§19](#).

MEGA Academy & Certificates

- `Course` → `CourseModule` → `Lesson`; `CourseEnrollment`; `Certificate`; `Instructor` (deliberately decoupled from `User` — can be named with no MEGA ID). See [§23](#).

Identity layer, content, and other supporting models

- `Interest` (self-declared), `Skill` (teacher-attested, `@@unique([studentId, addedByUserId, name])`).
- `Program`, `NewsPost`, `Resource`, `Event`, `Opportunity` — simple content types, no approval workflow. `Resource` and `Event` carry both a nullable `schoolId` and a nullable `organizationId` (polymorphic ownership); `Opportunity` the same. All three now have real write paths on both the School and Organization side: `Opportunity` gained Organization/School PATCH/DELETE (hard delete — no reverse relations) in K3; `Event/Resource` gained their Organization-side write paths in K8, `Event` reusing School Event's exact `isActive-deactivate` convention (no DELETE), `Resource` using hard DELETE (also no reverse relations, and the first write path it has had for either owner type).
- `Subscription/Payment` ■ — fully modeled, **no payment processor is connected**; priced-course enrollment is explicitly blocked.
- `Notification` — written only through `notify()/notifySchoolCommunity()`, never directly.

Cross-cutting schema conventions

- **Snapshot fields for anything that must survive a later rename:** `Certificate.*NameSnapshot`, `GradeHistoryAudit.previous*/new*` are plain values, not live FK relations. Logos are the one deliberate exception (looked up live).
- **The nullable-scope-discriminator idiom**, used identically across five different models: `TeacherAcademicAssignment.sectionId`, `TeachingUnit.sectionId`, `StudentEvaluation.gradeSubjectId`, `ParentTeacherMeeting.gradeSubjectId`, `AssessmentFrameworkAssignment.gradeSubjectId` — `null` means "general/grade-wide/default", a real value means "specific/override".

- **The recurring `NULL ≠ NULL` unique-index gap:** SQL treats two `NULL` values in a unique index as non-colliding, so a `@@unique` constraint with a nullable column cannot by itself block two rows that are both `null` in that column. This gap has been found (twice, live) and pre-checked (three times, proactively) across `TeacherAcademicAssignment`, `ClassTeacherAssignment`, `StudentEvaluation`, `AssessmentFrameworkAssignment`, and `AssessmentComponent` — each has an explicit application-level `findFirst`-based pre-check function rather than relying on the database constraint alone.
- **Most models have no working delete route** — for those, "delete behavior" is schema-level cascade behavior that would apply if a delete ever happened, never something a UI action triggers. The deliberate exceptions (real `DELETE` routes, because they're current-state operational data, not historical record): `TeacherGradeAssignment`, `GradeSubject`, `TeacherAcademicAssignment`, `ClassTeacherAssignment`, `AssessmentPeriod`, `AssessmentComponent` (until results exist), `AssessmentFrameworkAssignment` (until results exist).

(Source: [DATABASE.md](#))

8. Authentication & Authorization

Login

`src/lib/auth.ts` — NextAuth v4 with a single `CredentialsProvider` ("MEGA ID"): email + password, `bcrypt.compare` against `User.passwordHash`. `PrismaAdapter` is attached, but the session strategy is **JWT**, not database sessions. Custom sign-in page at `/login`.

Session handling

Every server-side check reads the session identically:

```
const session = await getServerSession(authOptions); const userId = (session?.user as any)?.id; const roles = (session?.user as any)?.roles;
```

The `jwt` callback copies `id/roles` onto the token at sign-in; the `session` callback copies them back onto `session.user`.

The `requireX()` helper suite (`src/lib/authorize.ts`)

Every write route that needs "does this session have permission to act on this specific resource" calls one of these rather than inlining the check. Each returns the authorized `userId` or `null` (never throws).

Function	Checks	Used for
<code>requireSchoolAdmin(schoolId)</code>	Session + a <code>SchoolAdmin</code> row	Nearly every <code>/api/schools/[id]/*</code> write route
<code>requireOrgAdmin(organizationId)</code>	Session + an <code>OrganizationAdmin</code> row	Every <code>/api/organizations/[id]/*</code> write route
<code>requireCourseOwner(courseId)</code>	Resolves course → org → org-admin check	Course content routes

<code>requirePlatformAdmin()</code>	Session + <code>roles.includes("PLATFORM_ADMIN")</code>	<code>/admin/*, /api/admin/*</code>
<code>requireSchoolFinance(schoolId)</code>	Session + (SchoolAdmin or SchoolAccountant)	Finance-only surfaces
<code>requireOrgFinance(organizationId)</code>	Session + (OrganizationAdmin or OrganizationAccountant)	Finance-only surfaces
<code>requireTeacherAssignment(schoolId, scope)</code>	A Teacher with an ACTIVE TeacherSchoolAffiliation at the requested school + a matching TeacherAcademicAssignment	Teaching Plans/Units/Tests, Subject Evaluations, Assessment Results marks entry/publish, Homework creation/edit (§17a)
<code>requireClassTeacher(schoolId, scope)</code>	A Teacher with an ACTIVE TeacherSchoolAffiliation at the requested school + a matching ClassTeacherAssignment	Attendance, General Evaluations
<code>teacherHoldsSubjectAssignment()</code> / <code>teacherHoldsClassAssignment()</code>	A named teacher's assignment (not the session's own user)	Admin-attributed evaluation/meeting creation

Deliberate design note: the finance helpers check *both* the Admin and Accountant relationships on purpose — an Admin keeps full authority (finance included); a bare Accountant gets finance access *only*.

The three-way `sectionId` scope semantics

`requireTeacherAssignment()`/`requireClassTeacher()` both accept an optional `sectionId` scope with **three distinct states**, resolved via a shared `sectionScopewhere()` helper:

- **omitted** — no section restriction, matches any assignment for the grade/subject.
- **null** — the target itself is grade-wide (e.g. a grade-wide TeachingUnit); requires a grade-wide assignment specifically.
- **a real section id** — the target is one specific section; a grade-wide assignment (covers every section) or that exact section's assignment both satisfy it.

This distinction was corrected in Phase 3B, before either helper had a real caller — the original implementation collapsed `null` and `omitted` into one code path (both are falsy in JavaScript), which would have wrongly authorized a section-specific-only teacher to manage a grade-wide unit.

Identity resolution — `resolveActiveTeacherId()` (Phase 4B)

Both `requireTeacherAssignment()` and `requireClassTeacher()` resolve the caller's Teacher identity through a shared helper, `resolveActiveTeacherId(userId, schoolId)`, rather than reading `Teacher.schoolId/Teacher.approved` directly: it looks up the caller's Teacher row, then requires an **ACTIVE** TeacherSchoolAffiliation for `(teacherId, schoolId)` — falling back to the legacy bridge-field check only when no affiliation row exists at all (pre-Phase-3 data never migrated). This is what allows a teacher with 2+ simultaneous ACTIVE affiliations to authorize correctly at every school they actually teach at, not just whichever single school `Teacher.schoolId` happens to hold. See [§10a](#) for the full affiliation model.

Phase 4C note: eleven other ownership/existence checks across the app (teacher/student address, contacts, contacts-address, skills, evaluations, the roster GET) previously re-derived identity a second time via the bridge fields — in two cases, *after* one of the helpers above had already authorized correctly through the affiliation table,

incorrectly rejecting an already-approved multi-school teacher. All eleven now resolve identity from `userId` alone, with no redundant bridge-based re-check.

Access patterns not covered by `authorize.ts`

Some checks are simple/specific enough to stay inlined: certificate preview access (`recipientUserId === userId || PLATFORM_ADMIN`), course enrollment/completion ownership, and the Promotion roster's closed-session access. The Student Profile page's own access check, `resolveStudentViewAccess()`, is a purpose-built resolver (`src/lib/institutionalContext.ts`) rather than an inline check — School Admin school-wide, or a Teacher whose `TeacherAcademicAssignment/ClassTeacherAssignment` actually covers the student's current section/grade, composing `requireTeacherAssignment()/requireClassTeacher()` above rather than the school-wide pattern other staff surfaces still use.

(Source: [AUTHENTICATION_AND_AUTHORIZATION.md](#))

9. User Roles & Permissions

All seven roles are stored identically as `UserRole` rows; a single `User` can hold **multiple roles at once**. `dashboard/page.tsx` picks which dashboard to render using a fixed priority order, explicitly commented in the code as *"a simple MVP priority order, not a permission hierarchy"*:

PLATFORM_ADMIN → SCHOOL_ADMIN → TEACHER → STUDENT → PARENT → ORGANIZATION_ADMIN → ACCOUNTANT

Role	Core capabilities	Approval
PLATFORM_ADMIN	Command-center dashboard (real counts, no invented metrics); verifies schools/organizations; can view any certificate. ■■ Can only be granted via <code>seed.ts</code> or direct DB access — no in-app route.	Seeded fixture only
SCHOOL_ADMIN	Full school management: profile, staff/student approval, Phase 2 (sessions/grades/sections/promotion), Phase 3A (subjects/teacher assignments), Phase 3B (attendance/units/tests), Phase 3C (evaluations/meetings), Phase 3D (assessment config/results), direct Student/Teacher creation.	Self-registers as first admin
TEACHER	Skill-crediting; course enrollment; within assignment scope — teaching plans/units/tests, evaluations, attendance (if also a Grade Coordinator/Class Teacher), assessment marks entry/publish.	School Admin (skipped if created via + Add Teacher)

STUDENT	Read-only: Teaching Progress, Test Results, Recent Attendance, shared Evaluations, published Assessment Results, Report Card, course enrollment/certificates.	School Admin (skipped if created via + Add Student)
PARENT	Read-only, per linked child: the same progress views as Student, plus Parent-Teacher Meetings (Student never sees these).	N/A (self-links children)
ORGANIZATION_ADMIN	Course authoring/publishing, opportunities, accountant management. ■■■ Organization.verified is never checked before publishing or enrollment.	Platform Admin (verified flag)
ACCOUNTANT	Finance-surface visibility only (no real transaction data exists — payments aren't integrated). Not self-registerable.	Granted directly by an Admin

Phase 4D addition: for a person connected to more than one school, the SCHOOL_ADMIN and TEACHER branches of `dashboard/page.tsx` (and the migrated feature pages below them) resolve *which* school to show through the institutional-context layer (`getAccessibleSchools()`/`verifySchoolAccess()`, §10a) before running their existing role-scoped queries — this sits in front of, and is independent of, the role priority-order above.

(Source: [USER_ROLES.md](#), [MEGA_ID.md](#))

10. School & Multi-School Architecture

A **MEGA ID belongs to the individual**, not to a school — a `User` is never owned by an institution. A person's relationship to a specific school is a row in `TeacherSchoolAffiliation/StudentSchoolAffiliation` (PENDING/ACTIVE/ENDED, time-bounded), not a single field on the account or even on the `Teacher/Student` profile — a person can hold zero, one, or several such rows simultaneously. This means:

- **Learning history follows the person, not the school.** `CourseEnrollment` and `GradeHistory` both link to `Student.id` → `User.id`; if a student transfers or leaves, their historical rows are never deleted — only a new status (TRANSFERRED/LEFT) is recorded.
- **A school and an organization are structurally independent.** A `Course` belongs only to an `Organization`; Phase 2's grade structure belongs entirely to a `School`. The two only meet at `Certificate.associatedSchoolId`, an informational link, never an ownership relation.
- **Registration is deliberately two-stage.** Every path creates a minimal, single-role account first; school/organization affiliation happens afterward from the person's own dashboard, via the JOIN/LEAVE/TRANSFER lifecycle described in §10a.
- **Multi-school independence:** nothing in the schema assumes a single school's data ever interacts with another school's. `AcademicSession`, `SchoolGrade`, `Section`, `Subject` are all `schoolId`-scoped, and every write route resolves the caller's own school via `requireSchoolAdmin/requireTeacherAssignment/requireClassTeacher` before touching any row.
- **A person can genuinely be affiliated with more than one school at once.** Teacher multi-school is explicitly designed for and tested (Phase 4B/4C authorize correctly against whichever school an ACTIVE affiliation

names). Student simultaneous multi-school is schema-permitted but not a decided product policy — see [§37](#).

(Source: [MEGA_ID.md](#), [ARCHITECTURE.md](#))

10a. Institutional Affiliation & Context Architecture

(Phase 3–4D. Source: [INSTITUTIONAL_CONTEXT.md](#), [MEGA_ID.md](#), [AUTHENTICATION_AND_AUTHORIZATION.md](#))

The three-layer model

User (MEGA ID, never school-scoped) → Teacher/Student (a stable role identity, still not itself school-scoped) → TeacherSchoolAffiliation/StudentSchoolAffiliation (the actual institutional relationship: status PENDING/ACTIVE/ENDED, real startDate/endDate, one row per person-school pair). A person can hold several affiliation rows, in any mix of statuses, at once.

Lifecycle: JOIN / LEAVE / TRANSFER / REJOIN

Implemented as primitives in `src/lib/affiliation.ts` (Phase 3), called from `POST /api/{teacher,student}/{join-school,leave-school,transfer-school}`:

- **JOIN** creates a PENDING affiliation; a School Admin approval (`POST /api/schools/[id]/[teachers,students]/[id]/approve`) flips it ACTIVE (resolved via the affiliation table first, bridge-field fallback only if no affiliation row exists at all).
- **LEAVE** ends the current ACTIVE affiliation (ENDED, endDate set); other affiliations at other schools are untouched.
- **TRANSFER** ends the old affiliation and creates a new PENDING one atomically. The primitives **throw** `AffiliationError` rather than returning `{ error }` — a correction made after discovering Prisma's `$transaction` only rolls back on a thrown error, never a returned error value, which had allowed a failed TRANSFER to leave a half-applied state.
- **REJOIN** is simply another JOIN after a prior affiliation at that school has ended.

Bridge fields — transitional, not authoritative

`Teacher.schoolId/approved/position/subjects` and `Student.schoolId/approved` still exist and are read by a shrinking set of not-yet-migrated call sites (see [§37](#)). They're kept in sync automatically only for the 0-or-1-open-affiliation case; once a person has 2+, the bridge fields are deliberately left untouched — never guessed — since there's no single correct value once the real relationship is one-to-many.

From ACTIVE affiliation to authorized access — `institutionalContext.ts`

```
getAccessibleSchools(userId: string): Promise<{ schoolId, schoolName, role: "SCHOOL_ADMIN" | "TEACHER" }[]> verifySchoolAccess(userId: string, schoolId: string): Promise<{ role: "SCHOOL_ADMIN" } | { role: "TEACHER"; teacherId: string } | null>
```

`getAccessibleSchools()` lists only **ACTIVE** relationships (PENDING excluded — an unapproved JOIN isn't a real relationship yet) and is a display/routing input only, never itself an authorization decision. `verifySchoolAccess()` is the actual security gate, re-checked fresh on every request with no caching — `null` (fail closed) for PENDING/ENDED/no relationship at all with that school. Scoped to School Admin and Teacher only; Student is deliberately excluded, since simultaneous multi-school policy for Student remains undecided ([§37](#)). Design

principle set explicitly for this and every migrated feature: ACTIVE affiliations alone define what's selectable, and the resolver is authoritative even in the single-school case — there's no separate "skip the resolver" path when a person has only one school.

The preference cookie — never authoritative

`mega_school_ctx` (POST `/api/dashboard/school-context`, `httpOnly`, `sameSite=lax`, 180-day) records a person's last-chosen school purely so a same-URL page doesn't force a re-choice every visit. Every read is followed by re-validating the school against `getAccessibleSchools()/verifySchoolAccess()`; a stale cookie pointing at a school the person no longer has ACTIVE access to is ignored, never trusted as a grant.

Three proven migration patterns

1. **URL-scoped** (`/dashboard/schools/[schoolId]/{attendance,evaluations,meetings}`) — the school is a path segment, `verifySchoolAccess()` checks it directly; a chooser (2+ schools) routes into a fresh URL per school. Established Phase 4D-1 (Attendance), reused unchanged Phase 4D-2 (Evaluations), Phase 4D-3 (Meetings).
2. **Same-URL** (`/dashboard/grades`) — the page's URL never changes; `SchoolChooser`'s `redirectTo` prop returns the person to that same URL, resolved via the freshly-set (and re-verified) cookie. Established Phase 4D-4 (Grades index).
3. **Target-derived** (`/dashboard/academics/[gradeSubjectId]`, `/dashboard/grades/[schoolGradeId]`) — no chooser: the URL parameter itself unambiguously belongs to one school, so the school is derived directly from the target entity, then checked with an exact `schoolAdmin.findUnique({ userId_schoolId })` — never an arbitrary `findFirst()` that could 404 a genuinely-authorized multi-school Admin. The Grades companion page was fixed to this pattern in Phase 4D-4, matching the `academics/[gradeSubjectId]` precedent.

Scoped client-side navigation — the `basePath` pattern

`AttendanceClient.tsx/EvaluationsClient.tsx/MeetingsClient.tsx` each perform in-page filter navigation (`router.push`). Before Phase 4D-3, all three hardcoded their unscoped legacy path, silently bouncing a person viewing the URL-scoped page back to the unscoped one on the first filter change. Fixed with an optional `basePath` prop on all three — supplied by the URL-scoped page to keep navigation scoped; omitted (the default) preserves the original unscoped behavior exactly.

What's still on the legacy pattern

Initial Setup, New Session, Assessment Frameworks, and Assessment Results still resolve school context via a plain `findFirst()` pick — in-progress migration debt, not a security gap (every write route independently re-checks ownership of the resource being changed). See §37. `/dashboard/profile` is no longer on this list (see §24) — it doesn't need a chooser at all, since it shows every affiliation/administered-school row for the caller at once rather than resolving to one "current" school.

Where this surfaces in the product — My Profile (My Profile K1)

`/dashboard/profile`'s **My Institutional Relationships** section is the first UI in the product to read `TeacherSchoolAffiliation/StudentSchoolAffiliation/SchoolAdmin` directly and present them to the person they're about — every relationship, including ENDED history, with dates rendered honestly (`startDateSource: "UNKNOWN_MIGRATED"` shows as "Not recorded," never a fabricated date). This is a pure read/presentation layer: it introduces no new query pattern beyond a `prisma.user.findUnique` with the relevant relations included, and does not touch `src/lib/affiliation.ts` or `institutionalContext.ts`. See §24.

Organization Institutional Context

(Source: [ORGANIZATION INSTITUTIONAL CONTEXT.md](#))

Organization's own, parallel foundation, following the exact `getAccessibleSchools()/verifySchoolAccess()` shape above:

```
getAccessibleOrganizations(userId: string): Promise<{ organizationId, organizationName,
role: "ORGANIZATION_ADMIN" | "ORGANIZATION_ACCOUNTANT" }[]> verifyOrgAccess(userId: string,
organizationId: string): Promise<{ role: "ORGANIZATION_ADMIN" } | { role:
"ORGANIZATION_ACCOUNTANT" } | null>
```

`getAccessibleOrganizations()` unions `OrganizationAdmin` with `OrganizationAccountant` links (display/routing input only). `verifyOrgAccess()` is a fresh, fail-closed gate, not yet called by any route — foundation for future organization-scoped gates. **Deliberately no schema change:**

`OrganizationAdmin/OrganizationAccountant` remain flat join rows with no status/history/dates — evaluated and explicitly rejected as a generalized `OrganizationMembership` table (no demonstrated need; School's own admin/accountant tables never got a history upgrade either — only the Teacher/Student *affiliation* layer did, because that gates real academic authority). `dashboard/page.tsx`'s `ORGANIZATION_ADMIN` branch now resolves via `getAccessibleOrganizations()` instead of an arbitrary `organizationAdmin.findFirst({ userId })`: 0 accessible → `CreateOrgPrompt`; 1 → `OrgDashboard`; 2+ → a non-interactive boundary page listing every accessible organization by name (deliberately **not** a chooser — no `OrganizationChooser`, no preference cookie, no URL-scoped route; that UX remains a separate, not-yet-approved decision).

Every existing role/resource-specific helper (`requireOrgAdmin`, `requireCourseOwner`, `requireOrgFinance`) is unchanged by this foundation and remains the actual authorization gate for every write route — `getAccessibleOrganizations()/verifyOrgAccess()` are a display/context layer, never themselves a security decision, exactly as `getAccessibleSchools()` is for School.

Organization logo, Events, and Resources (K6, K8) — all `requireOrgAdmin`-gated, all reusing this same institutional-context foundation for "who is this Organization's admin," none of them touching `getAccessibleOrganizations()/verifyOrgAccess()` directly (they use the resource-specific `requireOrgAdmin` gate, same as every other Organization write route). Full detail: [§23](#) (Academy participation distinction), [§24](#) (logo), [§26](#) (Events/Resources routes).

Organization institutional context vs. MEGA Academy participation — two independent facts. This section governs *who* administers an Organization and *what* it may do institutionally (post Opportunities/Events/Resources, once *verified*). `Organization.academyParticipant` (§23) governs *only* whether its MEGA Academy courses are publishable/enrollable/visible — a verified, fully-administered Organization may have `academyParticipant: false` and still operate normally otherwise.

11. Academic Sessions, Grades, Sections & Subjects

Academic Sessions

An `AcademicSession` represents one school year for one school — the container everything else in the academic system is scoped to. A school may have **at most one ACTIVE session** at a time, enforced at the application level (SQLite cannot express a partial unique index). Creating a session when one is already active returns `alreadyActive: true` (HTTP 200, not an error) rather than a second row — verified live against a genuine two-tab race condition.

Rollover (POST `/api/schools/[id]/academic-sessions/rollover`, `/dashboard/sessions/new`) closes the current session and opens a new one in one transaction: every student whose most recent decision was `COMPLETED/REPEATED` *with a real outcome grade* is auto-placed into the new session

(`carryForwardEligibleStudents()`); anyone still `ENROLLED` with no decision is left unplaced and surfaces in the persistent Pending/Unresolved queue. **Section is never carried forward** — the new row's `sectionId` is always `null`, regardless of the student's prior section.

Grade decisions (Promotion) happen **independently of session status** — a School Admin can promote Grade 6 today and leave Grade 7 untouched for weeks; nothing forces every grade to be resolved before the next rollover.

Grades

`SchoolGrade` is a school's opt-in to one platform-wide `GradeReference` (the fixed 13-rung ladder `PP1–PP3`, `Y1–Y10`), with its own `displayName` (e.g. "Class 6").

Sections

An **optional** subdivision of a `SchoolGrade` (Class 6 → A, B, C), added after the original Phase 2 build, on the same audit architecture:

- Belongs to the `SchoolGrade`, not to a session — the same rows persist across every academic session that grade is used in.
- Unlimited sections per grade; unique name per grade; **no hard-delete path** — deactivate (`isActive: false`) only, so historical `GradeHistory/GradeHistoryAudit` references always stay resolvable.
- School-Admin-only for every section-related action.
- **Never inherited across sessions or decisions** — a promotion, a repeat, and the rollover carry-forward sweep all create/leave the row's `sectionId` untouched from their own perspective (a decision never touches it) or explicitly `null` (a new session's row never inherits it). Section is always a separate, explicit action per session, mirroring how teacher assignments are never carried forward either.

Subjects (Phase 3A)

- **Subject** — a school-wide catalog entry, reusable across every grade and session; deactivate only, no hard delete.
- **GradeSubject** — which subjects a grade offers, **for one specific academic session** — every session starts with zero offerings for every grade; nothing is auto-copied from the prior session, so a past session's curriculum stays permanently reconstructable. A real `DELETE` route exists (blocked with 409 if a `TeacherAcademicAssignment` still references it).

(Source: [ACADEMIC_SESSIONS.md](#), [GRADES_AND_PROMOTION.md](#), [ACADEMIC_STRUCTURE.md](#))

12. Student Lifecycle & GradeHistory

`GradeHistory` is the central, permanent record of a student's grade placement for one session (`@@unique([studentId, academicSessionId])`). It is never deleted, and its `status/outcomeGradeId` may **only ever** be written through `recordGradeDecision()` (`src/lib/gradeHistory.ts`) — the sole audited write path. `sectionId` on an *existing* row may only ever be written through the separate `reassignSection()` function.

Lifecycle states

```
(no row) --create--> ENROLLED --decision--> COMPLETED | REPEATED | TRANSFERRED | LEFT
```

- **ENROLLED** — the default at creation; the only status eligible for a *new* Promotion decision.

- **COMPLETED** — promoted; `outcomeGradeId` required (the target grade).
- **REPEATED** — staying in the same grade next session; `outcomeGradeId` required (same grade).
- **TRANSFERRED** — left for another school; no outcome grade.
- **LEFT** — left the school entirely; no outcome grade.

First-time placement — four entry points, one shared architecture

A student's *first* `GradeHistory` row for a session can be created four ways, all sharing the identical "creation isn't a decision" rule (`direct GradeHistory.create()`, `status: "ENROLLED"`, no `decidedAt`, never routed through `recordGradeDecision()`):

1. **Initial Setup step 6** — `matchLegacyGradeText()` confident-match + manual queue, bulk, via `POST /api/schools/[id]/grade-placements`.
2. **Pending/Unresolved queue's "manually place"** — bulk, same route reused.
3. **Add Student, optionally, at creation time** — an inline `GradeHistory.create()` inside `POST /api/schools/[id]/students`'s own transaction.
4. **Students tab's "Assign Grade & Section →"** — a single-item call to the same `grade-placements` route, for any approved student with no `GradeHistory` row in the active session.

CURRENT_ROSTER_STATUSES — who counts as "currently in a grade"

`CURRENT_ROSTER_STATUSES (src/lib/gradeHistory.ts) = ["ENROLLED", "COMPLETED", "REPEATED"]` — the shared definition of "physically in this grade for the rest of this session," used identically by the Grades index's "N enrolled" count and the Class Overview roster. `COMPLETED/REPEATED` are included because a decision about *next* session doesn't remove a student from *this* grade the moment it's recorded — the school year isn't over. `TRANSFERRED/LEFT` remain excluded, since those students have genuinely left. This broadened definition **does not** change Promotion-panel eligibility — see [§22](#).

The legacy matching utility

`matchLegacyGradeText()` (`src/lib/gradeMatching.ts`) matches free text ("Grade 6", "VI", "Nursery") to a `GradeReference` code, handling the full Roman-numeral range with subtractive notation, ordinal suffixes, and pre-primary keywords. **Returns `null` — never a guess — for anything ambiguous**, verified with 20+ real inputs.

(Source: [GRADES AND PROMOTION.md](#), [DATABASE.md](#))

13. Teacher Assignments & Authorization Scopes

Three distinct assignment models, each with a different overlap policy:

Model	Scope	Overlap policy	Carried forward?
<code>TeacherGradeAssignment</code>	teacher → grade → session	N/A (no subject/section)	No

TeacherAcademicAssignment	teacher → grade → optional section → subject → session	Exclusive: a teacher may not hold both a grade-wide and section-specific row for the same subject/grade/session	No
ClassTeacherAssignment	teacher → grade → optional section → session (Grade Coordinator / Class Teacher)	Coexisting: grade-wide and section-specific rows may both exist for the same grade	No

TeacherAcademicAssignment.gradeSubjectId is a direct FK to the matching GradeSubject row — it is schema-impossible to assign a teacher to a subject the grade doesn't actually offer that session. Multiple different teachers may freely overlap on the same subject/grade/session (no teaching hierarchy exists or is planned).

requireTeacherAssignment() was built in Phase 3A *ahead of its first caller*, specifically as shared foundation for later phases — Phase 3B's Teaching Plans/Units/Tests, Phase 3C's Subject Evaluations, Phase 3D-2's marks entry/publish, and Phase 1 Homework creation/edit (§17a) all depend on it, exactly as the doc comment above it in authorize.ts anticipated. See §8 for its exact scope semantics.

(Source: [ACADEMIC_STRUCTURE.md](#), [ACADEMIC_OPERATIONS.md](#))

14. Attendance System

Attendance — one row per student **per calendar day**, never subject-based, @@unique([studentId, date]) (global, not per-session). date is always derived from an explicit client-sent "YYYY-MM-DD" string, never the server's own clock, so the value represents the school's intended calendar day regardless of server timezone.

Marking: POST /api/schools/[id]/attendance, requireSchoolAdmin OR requireClassTeacher scoped to sectionId (omitted = whole grade, requires a Grade Coordinator). A student whose actual placement doesn't match the marking pass's target grade/section is silently skipped; an already-marked student for that date is skipped, not an error.

Corrections go only through correctAttendance() (src/lib/attendance.ts), which updates the row **and** inserts an AttendanceAudit row capturing previousStatus/newStatus/previousRemarks/newRemarks together — every time, even a remarks-only edit records status unchanged, and vice versa.

Visibility: Student sees their own last-15-days attendance (read-only); Parent sees the same, per linked child; both reflect corrections, never the original pre-correction value.

Phase 4D-1: also reachable at the URL-scoped /dashboard/schools/[schoolId]/attendance, the proof-of-concept for the institutional-context migration pattern — same AttendanceClient, resolved and access-checked via verifySchoolAccess() instead of an unscoped bridge-field pick. See §10a.

(Source: [ACADEMIC_OPERATIONS.md](#))

15. Evaluations

`StudentEvaluation` — a teacher's narrative, qualitative remark about a student, for one session. **General vs. Subject is not a type field — it is entirely the presence or absence of `gradeSubjectId`:**

- `gradeSubjectId: null` — **General Student Evaluation**, authored by a Grade Coordinator / Class Teacher (`requireClassTeacher()`), on `/dashboard/evaluations`.
- `gradeSubjectId: <id>` — **Subject Evaluation**, authored by a Subject Teacher (`requireTeacherAssignment()`), on `/dashboard/academics/[gradeSubjectId]`.

A School Admin may create either **on behalf of a named teacher** — the named `teacherId` is independently validated (`teacherHoldsSubjectAssignment()/teacherHoldsClassAssignment()`), never simply trusted because the caller is an admin.

Audit-on-share

`updateEvaluationRemarks()` is the only code path that changes an existing evaluation's `remarks`:

- **While fully private** (`visibleToParent` and `visibleToStudent` both false) — a plain update, no audit row.
- **Once shared with EITHER audience** — every subsequent edit updates the row and inserts a `StudentEvaluationAudit` row, in the same transaction.

`visibleToParent/visibleToStudent` are **two fully independent gates** — sharing with one never implies the other; each has its own timestamp. Sharing itself is **one-way** — no un-share path exists in this phase.

The "STAFF" audience

Added for the Student Profile page (`/dashboard/students/[studentId]`) — skips the visibility filter entirely, returning every evaluation including private ones, since School Admins/approved Teachers already have full write access at their school.

Phase 4D-2: also reachable at the URL-scoped `/dashboard/schools/[schoolId]/evaluations`, reusing `EvaluationsClient` unchanged, resolved via the same institutional-context pattern as Attendance. See [§10a](#).

(Source: [ASSESSMENT AND EVALUATION.md](#))

16. Parent–Teacher Meetings

`ParentTeacherMeeting` handles both **occasional** (single) and **periodic** (bulk) meetings through one model and one route (`POST /api/schools/[id]/meetings, {meetings: [...]}` — one item vs. many). No recurrence/series concept exists.

- **Initiation:** School Admin or an authorized Teacher only. **Parents have no write path at all** — read-only recipients in this phase.
- **Every item is resolved/validated before the transaction opens** — the Postgres-safe pattern, not the SQLite-only catch-mid-transaction pattern used by `grade-placements/teacher-assignments`.
- **Editing** (`PATCH .../meetings/[meetingId]`) is authorized by identity — the meeting's own (still-approved) teacher, or a School Admin.
- **Rescheduling** (`scheduledAt/location/onlineUrl`) is only allowed while `status: "SCHEDULED"` — a `COMPLETED/CANCELLED` meeting's details are historical record.
- **linkedEvaluationId** — a plain, non-unique FK to a `StudentEvaluation`, for attaching prepared context.

Structural, not just UI, visibility guarantee

`fetchMeetingsForStudent(studentId, audience)` (`src/lib/academicProgress.ts`) accepts **only** "PARENT" | "STAFF" — its type signature has no "STUDENT" member at all. It is called only from the Parent dashboard branch and the staff-only Student Profile page — never from the Student branch, and never folded into the shared `AcademicProgressPanel.tsx` that the Student's own render path uses. A Student attempting direct navigation to `/dashboard/students/[studentId]` or `/dashboard/meetings` is redirected away before any meeting data is fetched.

Phase 4D-3: also reachable at the URL-scoped `/dashboard/schools/[schoolId]/meetings`, reusing `MeetingsClient` unchanged, resolved via the same institutional-context pattern. This phase's own audit found the Meetings API routes (POST `/api/schools/[id]/meetings`, PATCH `.../meetings/[meetingId]`) still resolving the calling teacher via the legacy bridge fields, never touched by Phase 4B/4C — fixed to use `resolveActiveTeacherId()` the same way `requireTeacherAssignment()/requireClassTeacher()` do. See [§10a](#).

(Source: [ASSESSMENT_AND_EVALUATION.md](#))

17. Unit Tests

`TeachingUnit` — one curriculum unit/chapter under a `GradeSubject`, with `status`: `NOT_STARTED` | `IN_PROGRESS` | `COMPLETED` (managing `startedAt/completedAt` automatically) and an app-assigned `order`. `sectionId`: `null` means a grade-wide unit sequence shared by every section; a real value means that section's own independent sequence.

`TeachingPlan` — a **separate model**, not a field on `TeachingUnit`: a teacher's declared `plannedTotal` (independent of how many `TeachingUnit` rows actually exist) and `unitLabel` (free text — "Unit" or "Chapter", a plain string, not a second model).

`UnitTest` — belongs to one `TeachingUnit`; **creatable only once the unit is `IN_PROGRESS` or `COMPLETED`**. `UnitTestResult` is **pre-created** (`status`: "PENDING") for every student currently enrolled in the test's scope at test-creation time — a stable roster snapshot, deliberately the opposite pattern from `AssessmentComponentResult`'s lazy creation (see [§19](#)) because a test has one clean, one-shot creation moment. `ABSENT` forces `marksObtained` to `null`; `EVALUATED` requires $0 \leq \text{marksObtained} \leq \text{maxMarks}$.

`UnitTest/UnitTestResult` remain **genuinely separate from `AssessmentFramework/AssessmentComponent`** — no merge, no auto-derivation of marks between the two systems, by explicit design decision.

(Source: [ACADEMIC_OPERATIONS.md](#), [ASSESSMENT_FRAMEWORK.md](#))

17a. Homework — Phase 1

`Homework` (`prisma/schema.prisma`) is a Teacher-authored item scoped to `academicSessionId/schoolGradeId/optional sectionId/gradeSubjectId/subjectId`, for one grade (or one section of it) and one subject, for one session. **No per-student fan-out row exists** — unlike `UnitTestResult`, one `Homework` row serves every applicable student; visibility is resolved entirely at read time against a student's own `GradeHistory` placement, reusing the same three-way `sectionScopeWhere()` idiom already established for teacher authorization (see [§8](#)).

- **Authorship:** `teacherId` → `Teacher`, not `createdByUserId` → `User` — a deliberate divergence from `TeachingUnit`'s precedent, matching `StudentEvaluation`'s instead, so authorship resolves through the real institutional Teacher identity, never the raw account.

- **Lifecycle:** a single `status` (DRAFT/PUBLISHED) plus `publishedAt` gates Student **and** Parent visibility together — not two independent flags like `StudentEvaluation.visibleToParent/visibleToStudent`, since nothing in Phase 1 calls for an asymmetric audience between a student and their own parent. Once PUBLISHED, `title/instructions/dueDate/sectionId` are frozen (PATCH returns 409), matching the "permanent once shared" precedent already used for `StudentEvaluation` sharing and Certificate issuance; only the DRAFT → PUBLISHED transition itself remains allowed, idempotently.
- **Date semantics:** `dueDate` is date-only, following the `Attendance.date/UnitTest.testDate` convention (client-supplied "YYYY-MM-DD", exact-match). "Today's Homework" means **dueDate equals today** — not "published today." "Today" is resolved via `todayInKathmandu()` (`src/lib/homework.ts`), an `Intl.DateTimeFormat` call hardcoded to `Asia/Kathmandu` — deliberately **not** the naive `new Date().toISOString().slice(0, 10)` pattern used elsewhere in this codebase (e.g. `dashboard/attendance/page.tsx`, purely as an overridable UI default), since Nepal Standard Time (UTC+5:45) makes that pattern show the wrong calendar date for roughly the first ~5h45m of every Nepal day. No date/timezone package was installed; no `School` timezone field was added — the anchor is hardcoded, since the platform is Nepal-only today.
- **Shared read function:** `fetchTodaysHomework(studentId)` is the one place this query is written, called by both the Student dashboard branch and, once per linked child, the Parent branch — mirroring `fetchAcademicProgress()`'s "one function, every caller" discipline (see §25, §30). Like `fetchAcademicProgress()`, it performs no authorization itself — callers are responsible for only ever passing a `studentId` they've already verified the caller may see.
- **Parent visibility:** reuses the existing `Parent → ParentStudent → Student` chain unchanged — **no new ParentHomework relationship** was introduced. `dashboard/page.tsx`'s PARENT branch adds `fetchTodaysHomework(c.student.id)` to its existing per-child `Promise.all` fan-out (alongside `fetchAcademicProgress/fetchMeetingsForStudent/fetchAssessmentResults`), with the identical security posture: child ids are derived entirely from the caller's own `ParentStudent` rows, never a request-supplied `studentId`.
- **Authorization:** `POST /api/schools/[id]/homework (create)` is gated by `requireTeacherAssignment()` alone — deliberately **not** composed with `requireSchoolAdmin()` (unlike `TeachingUnit`'s create route), since `Homework.teacherId` is a real, non-nullable `Teacher` FK and Phase 1 has no approved "Admin authors on behalf of a teacher" flow. `PATCH .../homework/[homeworkId] (edit/publish)` accepts either `requireSchoolAdmin()` or `requireTeacherAssignment()` re-checked fresh against the homework row's own stored scope — never just "are you the teacher who originally created this." Every relational id (`gradeSubjectId`, `sectionId`) and the homework row's own `schoolGradeId` are re-validated against the URL's `schoolId` server-side — defense in depth against a forged cross-school `homeworkId`, verified live.
- **UI:** `/dashboard/schools/[schoolId]/homework` — URL-scoped from day one via `verifySchoolAccess()` (§10a); unlike `Attendance/Evaluations/Meetings`, Homework has **no unscoped legacy sibling**, since it has no pre-Phase-4D history to preserve. Reachable from the School Admin and Teacher dashboards and the Phase 4D-1 multi-school chooser page, the same discoverability approach already used for the Public School Gateway's Inquiries.
- **Deliberately out of scope for Phase 1** (see §37): submissions, attachments, grading, rubrics, feedback, discussion, plagiarism checking, analytics, reminders/notifications on publish, upcoming/past views, and a School-Admin-authors-on-behalf-of-teacher flow.

(Source: [HOMEWORK.md](#))

17b. Calendar — Kilometer 1 / 1.1 / 1.2

(Source: [CALENDAR.md](#))

A **projection layer**, never a second source of truth — no `CalendarEvent` model exists. `CalendarItem` (`src/lib/calendar.ts`) is a plain TypeScript shape assembled by source-specific adapter functions from

already-authoritative models: `GeneralCalendarEntry` (new, network-wide reference dates), `Event` (pre-existing, previously had no write path at all — confirmed zero rows in every environment before Kilometer 1), `ParentTeacherMeeting`, `Homework`, `AcademicSession`, and `SchoolCalendarEntry` (new, Kilometer 1.1). Every item is traceable back to its real source row.

Day Status — an independent layer from Activities

Kilometer 1.1 added **Day Status** ("what kind of day is this?"), deliberately never merged into `CalendarItem` ("what happens on this day?"). `resolveDayStatuses(schoolId, window)` (`src/lib/schoolCalendar.ts`) resolves exactly one dominant value per date — `SPECIAL_CLOSURE` > `EXAMINATION` > `VACATION` > `PUBLIC_HOLIDAY` > `WEEKLY_HOLIDAY`, no blended/dual-color treatment — from `SchoolCalendarEntry` ranges, `GeneralCalendarEntry` rows where `type === "NATIONAL_HOLIDAY"` only (`OBSERVANCE/MEGA_WIDE_EVENT` never set a status — an observance never implies school is closed), and a computed weekly holiday (Saturday, hardcoded — no `School` field, no per-school configuration). A Parent whose children attend different schools gets each school's statuses resolved independently, then merged with the same priority (`mergeDayStatuses()`). Kilometer 1.2 gave each status a distinct color family (blue/emerald/amber-200/slate/rose) after the original palette (several pale near-identical neutrals) proved too subtle in a dense grid — visual-only, no change to the resolution logic or priority order.

`SchoolCalendarEntry` (Kilometer 1.1)

One new model covering six categories: `VACATION/EXAMINATION/SPECIAL_CLOSURE` (date ranges, `affectsDayStatus: true`) and `PTM/RESULT_DAY/REPORT_CARD_DISTRIBUTION` (single-date point activities, `affectsDayStatus: false`). `affectsDayStatus` is **server-derived from category alone** (`resolveSchoolCalendarEntryDates()`) and never accepted from the client, even if supplied — verified live: a request sending `{category: "PTM", affectsDayStatus: true}` is stored with `affectsDayStatus: false`. School-Admin-only write path (`POST/PATCH /api/schools/[id]/school-calendar[/entryId]`), soft-deactivate only, no `DELETE` — mirrors `Event`'s exact write-path shape. Deliberately not where individual `ParentTeacherMeeting` appointments, `UnitTest` dates, or `AssessmentPeriod` live.

Role-aware aggregation, authorization, and links

School Admin/Teacher reach Calendar via `/dashboard/schools/[schoolId]/calendar`, gated by `verifySchoolAccess()` ([§10a](#)) exactly like Attendance/Homework/Meetings. Student/Parent resolve their own scope from session identity (`/dashboard/calendar`, no `schoolId` in the URL). Parent aggregation carries two optional `CalendarItem` fields, `childId/childName`, populated only for genuinely per-student sources (Homework, individual `ParentTeacherMeeting`) — school-wide sources (Events, `SchoolCalendarEntry`, Academic Session boundaries) are fetched once per **distinct school**, never once per child, so two siblings at the same school never see that school's Events/Vacation/Exam dates doubled. (A real bug in this exact area — a shared Homework row producing two `CalendarItems` with a colliding id — was caught and fixed by qualifying the id with the child id when attributed.) Public `/calendar` shows General Calendar unconditionally plus a search-first, single-school-bounded institutional section — replacing the pre-K1 unscoped `prisma.event.findMany({take:50})`, which would have starved smaller schools out of visibility at scale. `CalendarItem.link` points only at real, already-existing destinations (Homework → the school's Homework page for Teacher/Admin; `ParentTeacherMeeting` → `/dashboard/meetings?teacher=<id>`; `RESULT_DAY/REPORT_CARD_DISTRIBUTION` → the existing `/dashboard/report-card/[studentId]` for Student/Parent) — never an invented page.

Views

Annual (default) — 12 static month cards (`CalendarAnnual.tsx`), `Date.UTC`-based grid math (`src/lib/monthGrid.ts`), Day Status background plus the existing per-item activity dot/number treatment layered on top, one always-visible legend. **Agenda** — the original chronological list (`CalendarAgenda.tsx`, unmodified

since Kilometer 1), kept to its own ~30-day upcoming window regardless of how wide the Annual view's shared fetch is — both views render from the same fetched `CalendarItem[]`, no second query. Neither is an interactive drag/drop scheduling grid.

Event management

School-Admin-only POST/PATCH `/api/schools/[id]/events[/eventId]` (Kilometer 1) — Create, Edit, and soft-Deactivate (`isActive:false`), no DELETE route. `CalendarEventForm.tsx` (Kilometer 1.1) exposes Edit/Deactivate for the school's own upcoming Events; `SchoolCalendarManager.tsx` provides the equivalent for `SchoolCalendarEntry`. Both are deliberately kept out of `CalendarAnnual/CalendarAgenda`, which stay display-only and shared by every role.

Deliberately out of scope

BS↔AD conversion of any kind (`GeneralCalendarEntry.bsDateDisplay` remains a plain, unpopulated display string — no library installed, no fabricated value), Grade/Section/Subject or category filtering on the Admin Calendar, per-school configurable weekly holiday days, blended Day Status treatment, recurring events, calendar sync/export/subscriptions, notifications, an interactive Week/Day scheduling grid, and Organization Calendar (Organizations have no institutional-context parity with Schools — see [§10a](#), [§37](#)). The seeded `GeneralCalendarEntry` reference data covers September–December 2026 only and rests on a secondary source (`qppstudio.net`, cross-referenced) — not yet confirmed against the Nepal Panchanga Nirnayak Bikash Samiti or Ministry of Home Affairs gazette, and never described as final.

18. Assessment Framework System

Real report cards from different schools use genuinely different marking structures — percentage-weighted with letter grades, raw maximum marks, term-split, or a different structure per subject. Phase 3D-1 built a **configurable** assessment system rather than one hard-coded structure.

Core models

- **AssessmentFramework** — a reusable, **school-wide** marking-scheme template, **not session-scoped** (like `GradingScale`, and like `Subject/GradeSubject`'s catalog-vs-offering split). Deactivate only, no hard delete.
- **AssessmentPeriod** — an optional grouping layer ("Term I", "Annual") — a framework may define zero periods (flat structure) or several.
- **AssessmentComponent** — one scored/graded/descriptive piece, belonging directly to a framework or to one of its periods. `maxMarks: Float` does **double duty as marks and weight** — a component worth 10% and a component worth 10 raw marks are the identical shape; only `entryMode` (`MARKS` | `GRADE` | `DESCRIPTIVE`) differs.
- **GradingScale/GradingScaleBand** — a reusable, school-wide marks→grade conversion table (`minPercent/maxPercent/label/gradePoint?/isPassing?`).
- **AssessmentFrameworkAssignment** — the **only session-scoped model** in this phase: binds a framework to (`AcademicSession`, `SchoolGrade`), optionally narrowed to one `GradeSubject` as a subject override.

Resolution rule

`resolveFrameworkAssignment()` (`src/lib/assessmentFramework.ts`): look up a subject-specific assignment first; if none exists, fall back to the grade-default. The same nullable-scope-discriminator idiom used four other places in this schema.

Setup UI — guided wizard as the primary path

`/dashboard/assessment-frameworks/new` — **"Create Assessment System"** is the primary, promoted way a School Admin sets up a framework: a 6-step guided flow in plain language (Name it → Terms or one overall assessment? → How are students assessed? → How should results be shown? → Where should this apply? → Review & confirm), deliberately hiding the underlying "framework"/"assignment" terminology. **Zero new API routes** — confirming the wizard sequences the exact same existing calls: an optional `POST /grading-scales`, then `POST /assessment-frameworks`, then `POST /assessment-framework-assignments`.

`/dashboard/assessment-frameworks` itself is the landing page: a plain-language "Your Assessment Systems" list with the wizard promoted above it, and the original detailed editor (grading scales, per-framework period/component editing, the raw assignment form) tucked behind a collapsed **"Advanced management"** disclosure for anyone editing something after creation.

Authorization

School-Admin-only for every write route in this phase — no teacher-facing action exists until marks entry (Phase 3D-2).

(Source: [ASSESSMENT_FRAMEWORK.md](#))

19. Assessment Results & Publishing Workflow

Lazy result creation

Unlike `UnitTestResult`'s eager pre-creation, `AssessmentComponentResult` rows are created **only when a value is actually written**. A component is defined at framework-*design* time — long before any assignment exists — and can be added to an already-assigned, reusable framework; eager pre-creation would silently miss any component added afterward. The marks-entry page instead computes a **virtual roster** (`GradeHistory`, scoped to the assignment's grade) and left-joins whatever result rows already exist.

Publishing workflow

(no row) → first entry → DRAFT → explicit Publish → PUBLISHED → correction → PUBLISHED + audit

- **Publication is subject-level** (`AssessmentResultPublication`, keyed by `gradeSubjectId` + `studentId`), **never per-component** — a Parent/Student can never see a partially-published subject.
- **Who can publish:** `requireSchoolAdmin(id) || requireTeacherAssignment(id, {..., subjectId})` — the identical composition already used for `StudentEvaluation` sharing. Being a Grade Coordinator or Class Teacher alone does not provide special authority to enter, correct, or publish assessment results — assessment authority requires the appropriate matching `TeacherAcademicAssignment` for the subject.
- **Completeness is computed server-side**, not trusted from the client: `POST .../subjects/[gradeSubjectId]/publish` silently skips (never errors) any student with a still-PENDING non-DESCRIPTIVE component.
- **A correction to an already-PUBLISHED result stays published and is audited** (`AssessmentComponentResultAudit`) — it is **never reverted to DRAFT**, matching the identical `StudentEvaluation` precedent.

ABSENT handling

status: "ABSENT" is settable regardless of entryMode, forces marksObtained/gradeLabel/remarks to null regardless of what's passed, and contributes zero (not null) to numeric aggregation.

Data Integrity Protections (three guards on existing 3D-1 routes)

Guard	Rule
Component structural lock	maxMarks/entryMode locked once any result exists; renaming stays free
Assignment deletion block	DELETE rejected once any result/publication references it
Grading-scale lock	Full bands replacement rejected once any PUBLISHED result uses the scale; name/isActive remain editable

(Source: [ASSESSMENT_RESULTS.md](#))

20. Grading & GPA Calculations

All calculation lives in src/lib/assessmentResults.ts — **nothing is cached**; every total, percentage, and grade is recalculated on every read.

Function	Purpose
computeComponentContribution()	Turns one raw result into a point value. MARKS → raw value; GRADE → the matched GradingScaleBand's percentage midpoint applied to maxMarks (deliberately not gradePoint, which is an arbitrary school-chosen scale with no inherent percentage relationship); DESCRIPTIVE → always excluded from every numeric total; ABSENT → zero; PENDING → null (never zero).
aggregateGroup()	One function for both period-scoped and framework-level aggregation — components → (periods) → subject, called identically at every level.
lookupGrade()	Matches a percentage to a GradingScaleBand, lower-inclusive/upper-exclusive (minPercent <= p < maxPercent) — a shared boundary (e.g. 80%) is never ambiguous; a percentage of exactly 100 matches the band whose maxPercent is 100. This boundary rule was corrected during verification after an inclusive-both-ends bug let a boundary score match two bands.
computeUnweightedGPA()	Plain arithmetic mean across whichever subjects have a resolvable gradePoint — no subject-credit/weighting concept anywhere in this schema. null (not 0) if no subject has one.

<code>computeUnweightedAveragePercentage()</code>	The fallback ranking basis when no subject resolves a GPA (e.g. a label-only grading scale with no <code>gradePoint</code> anywhere) — the same unweighted-average shape, averaging each subject's percentage instead. Added specifically for Class Overview ranking.
---	---

GPA/percentage are never forced where the data doesn't support them: a marks-only framework produces a percentage but no grade/point; a descriptive-only framework produces neither — both fall out naturally from the same functions, with no special-casing.

(Source: [ASSESSMENT_RESULTS.md](#))

21. Report Card Architecture

`buildReportCard(studentId, audience)` assembles student/school/session/grade info, `fetchAssessmentResults()`'s subject results and GPA, and — reused directly, **not re-queried** — `fetchAcademicProgress()`'s attendance and evaluations. Rendered live at `/dashboard/report-card/[studentId]`.

Deliberately not modeled like Certificate. A certificate is a permanent, frozen-at-issuance snapshot; a Report Card must reflect corrections made after publication, so freezing one into a stored row would directly contradict the audited-correction design. **There is no persisted ReportCard model and no PDF export** — both explicitly out of scope.

Access: the Student themselves, a linked Parent, or staff (School Admin / any approved Teacher at the school, no assignment-level scoping — the same Skills-page precedent; the Student Profile page itself no longer follows this pattern — see [§8](#)) — resolved to the matching audience before calling `buildReportCard()`, so a Student/Parent only ever sees published data through the identical filter used everywhere else.

(Source: [ASSESSMENT_RESULTS.md](#))

22. Promotion & Grade Decisions

`/dashboard/grades/[schoolGradeId]` — a School Admin opens one grade's Class Overview, multi-selects students from the roster, and applies one decision to the whole selected batch, entirely inside one `prisma.$transaction` (verified with a real 100-student batch: 328ms, decided: 100, skipped: 0).

Decision	status	outcomeGradeId
Promote	COMPLETED	required — the target grade
Repeat	REPEATED	required — same grade
Transfer	TRANSFERRED	none
Leave	LEFT	none

Eligibility is re-checked server-side: submitting an already-decided or foreign `gradeHistoryId` is silently excluded and counted in `skipped`, never double-decided. **Promotion is completely independent of section** — it never reads or writes `sectionId`; the roster's "Apply Decision" and "Assign Section" panels hit entirely separate endpoints.

Class Overview vs. Promotion-panel eligibility

The Class Overview roster uses the broadened `CURRENT_ROSTER_STATUSES` (see §12) so the page honestly shows everyone currently in the grade, decided or not. **The Promotion-action panel's own eligibility stays deliberately narrower** — only a still-`ENROLLED` row can be selected for a *new* decision, enforced independently server-side by the `grade-decisions` route regardless of what the roster displays. Broadening the display never changes who a decision can actually be applied to.

Class Overview additions

- **Teachers & Subjects** — every `TeacherAcademicAssignment` for the grade/session.
- **Section-wise grouping** — by the student's current-session `sectionId`; a final "Unassigned / No Section" group, never hidden; a per-section, **display-only** Roll No. (not a persisted field — see §37).
- **"Repeated" badge** — derived from the student's **prior-session** row only (`status === "REPEATED"` and its `outcomeGradeId` equals this grade), never from their current-session status — a forward-looking decision about *next* session is never conflated with how they arrived in *this* one.
- **Top 5 ranking** — computed **once, across the whole grade**, before any section grouping; reuses `fetchAssessmentResults()/computeUnweightedGPA()/computeUnweightedAveragePercentage()`, filtered to published results only.

Undecided students & the Pending/Unresolved queue

A student never decided stays `ENROLLED` indefinitely. On rollover, they're excluded from automatic placement and surface in a **persistent** Pending/Unresolved queue on `/dashboard/grades` — resolvable by recording the missing decision (audited) or manually placing them directly (unaudited, an honest gap in the record). Verified correct across a 3-session chain, including a student pending through an intervening session with zero rows of their own.

Phase 4D-4: institutional context migration

`/dashboard/grades` was migrated to the **same-URL** institutional-context pattern (a School Admin with 2+ `ACTIVE` affiliations sees a chooser, but stays on `/dashboard/grades` itself — resolved via the just-set `mega_school_ctx` cookie rather than a new URL). This surfaced a second, independent bug in the companion `/dashboard/grades/[schoolGradeId]` page: it resolved its own school via an arbitrary `schoolAdmin.findFirst({ userId })`, which would 404 a multi-school Admin clicking into a grade belonging to any school other than whichever one `findFirst()` happened to return. Fixed via the **target-derived** pattern already used by `academics/[gradeSubjectId]`: `schoolId` is read directly from the target `SchoolGrade` row, then checked with an exact `schoolAdmin.findUnique({ userId_schoolId })`. See §10a.

(Source: [GRADES_AND_PROMOTION.md](#))

23. Certificate System

MEGA Academy — Organization provider architecture, participation, and enrollment

(Source: [COURSES AND ENROLLMENTS.md](#), [ORGANIZATION INSTITUTIONAL CONTEXT.md](#))

Course (belongs to an Organization, optionally an EducationalApproach/Instructor) → CourseModule → Lesson, plus CourseEnrollment and its optional 1:1 Certificate.

- **User-centric enrollment:** CourseEnrollment.userId (required, direct User FK) is the enrollment's real identity — any authenticated MEGA ID may enroll, independent of institutional role. teacherId/studentId remain on the model as **optional contextual enrichment only** (populated when the enrolling user happens to also hold that profile, purely so TeacherDashboard/StudentDashboard's reverse-relation queries keep working) — never re-derived as the enrollment's identity, both null for a Parent/Organization Admin/Accountant/unaffiliated enrollee. @@unique([courseId, userId]).
- **Organization.academyParticipant** — a self-service boolean, independent of Organization.verified. Publishing a course (PATCH /api/courses/[courseId], published: true) and enrolling (POST /api/courses/[courseId]/enroll) both require **both** verified and academyParticipant to be true. Toggled by the Organization Admin (PATCH /api/organizations/[id] — accepts only this field), checked fresh on every request, never cached; turning it off blocks new publish/enroll attempts immediately and has zero effect on existing courses, enrollments, or certificates.
- **Global read-time visibility:** /courses, /courses/[slug], and the homepage's course query all key off the identical published && organization.verified && organization.academyParticipant && organization.isActive condition — a course disappears from every one of these surfaces the instant any one fact goes false, and reappears the instant it's true again, with the course row/enrollments/certificates never touched. /courses/[slug]/learn (an already-enrolled learner's own access) is deliberately exempt from this gate — historical access persists regardless of the organization's current state.
- **Public Organization provider profile** (/organizations/[slug]) — public, unauthenticated, re-verifies verified && isActive independently at the detail-page level (never relying only on /organizations's list filter, since a slug is guessable). Two independent badges: "✓ Verified Organization" (always, since the page requires it) and "✓ MEGA Academy Provider" (only if academyParticipant) — never combined into one compound status. Shows published courses (only when participating), Opportunities, Events, and Resources (all independent of academyParticipant — see below).
- **Course → Provider reciprocal links** — the organization name on /courses, /courses/[slug], and the homepage's course cards links to /organizations/[organization.slug], completing the reverse direction of the provider→course journey the profile page established.
- **Opportunity, Event, Resource ownership** — all three belong to Organization (or School) independently of academyParticipant; an Organization's general public presence (Opportunities/Events/Resources) is never gated by whether it currently offers MEGA Academy courses. Opportunity has full POST/PATCH/DELETE on both School and Organization sides (hard delete — no reverse relations). Event/Resource gained their Organization-side write paths in K8 — see [§10a](#) and [§26](#).

Certificate keeps **recipient, instructor, and issuer as separate concepts**, never conflated — recipientUserId, optional instructorId (an Instructor can be named with no MEGA ID at all), issuerType (MEGA_EDU | ORGANIZATION | SCHOOL | JOINT — only ORGANIZATION is reachable today), and a separate associatedSchoolId for informational school context. Certificate.recipientUserId already addressed recipients by User.id natively — the user-centric enrollment work above required no change to certificate issuance.

Issuance: issueCourseCertificate() (src/lib/certificates.ts) is the **only** code path that creates a Certificate for a course — called atomically alongside marking a CourseEnrollment complete, so an enrollment can never end up "complete" with no certificate or vice versa.

Grade certificates ■ — Certificate.gradeHistoryId is a reserved, unlinked column; issueGradeCertificate() **does not exist**. This was a deliberate Phase 2 scope exclusion, not an oversight — GradeHistory fully exists, but the issuance path was never built.

Display: `CertificateDocument.tsx`, a true-to-size A4 landscape layout driven by a pure view-model builder (`buildCertificateViewModel()`), reading only frozen snapshot fields plus a live logo lookup (the one deliberate exception to the snapshot rule — see §27). Shown at `/dashboard/certificates/[id]/preview` (access-gated to the recipient or a Platform Admin) and publicly, unstyled, at `/verify/[code]`.

Not built: PDF export ■, QR code generation ■ (a space is marked but empty). A school with a logo uploaded (see §24) shows it here. `Organization.logoUrl` now exists (K6 — see §24) but was **not** wired into certificate rendering by that kilometer — an organization-issued certificate still always renders the name-only fallback, explicitly unchanged, not a broken state.

(Source: [CERTIFICATES.md](#))

24. School Logos & User Profile Photos

An identity/media-upload layer added on top of the existing `School` and `User` models — no new models, one additive field on each.

Data model

- `School.logoUrl` (pre-existing field, previously unpopulated) — now actively managed. Displayed on School Directory cards, a school's public profile page, the homepage's "Recently joined schools," and the School Admin dashboard's identity header.
- `User.avatarUrl` (new field) — deliberately placed on the shared `User` model, **not** duplicated onto `Teacher`, `Student`, or `Parent`. This follows MEGA ID's core principle directly: **one identity, many roles** — a person's photo is a property of *them*, not of any one role they hold, so the same avatar automatically appears across every dashboard/list surface regardless of which role is being viewed.

Organization logo (K6)

`Organization.logoUrl` `String?` (nullable, db push-added — no migration file). `POST/DELETE /api/organizations/[id]/logo`, gated by `requireOrgAdmin(organizationId)` only — **not** `requireOrgFinance`, so an Accountant cannot set or remove it. Reuses `saveUploadedImage/deleteUploadedImage` from `src/lib/uploads.ts` completely unchanged (same validation, same replace/delete-ordering lifecycle described below), storing under `public/uploads/organizations/{id}/` alongside `public/uploads/schools/{id}/`. Displayed via the same shared `Avatar` component (`variant="school"` — the existing institutional-mark styling; no new variant was added) on the Organization Dashboard, `/organizations`, and `/organizations/[slug]`.

API routes & authorization

Route	Methods	Authorization
<code>/api/schools/[id]/logo</code>	POST (upload/replace), DELETE (remove)	<code>requireSchoolAdmin(schoolId)</code> — the same helper every other school-profile write route uses
<code>/api/organizations/[id]/logo</code>	POST (upload/replace), DELETE (remove)	<code>requireOrgAdmin(organizationId)</code> — mirrors the School route exactly

<code>/api/user/avatar</code>	GET (read own), POST (upload/replace), DELETE (remove)	Session-based self-service only — the caller's own <code>userId</code> , inline-checked the same way <code>/api/interests</code> already does. There is no route or code path through which one user can set another user's <code>avatarUrl</code> .
-------------------------------	--	--

Upload validation

Enforced server-side in `src/lib/uploads.ts` — the client-supplied filename and browser-reported MIME type are never trusted:

- Accepted formats: **PNG, JPEG, WebP** — detected by reading the file's actual magic bytes, not the filename extension or the `Content-Type` header. **SVG is not accepted** (avoids script-injection risk from uploaded SVG markup).
- Maximum size: **2MB**.
- Stored filenames are always a fresh server-generated UUID (`crypto.randomUUID()`) plus the sniffed extension — the original uploaded filename is never used or persisted anywhere.
- A file that fails validation (wrong format, too large, or spoofed type) is rejected before anything is written to disk — no partial or orphaned file results from a rejected upload.

Storage architecture

Local filesystem, under `public/uploads/{schools,users}/...` (gitignored — never committed; user-uploaded content is never part of the repository). `logoUrl/avatarUrl` store the resulting root-relative URL, served by Next.js's normal static file handling and rendered through `next/image` for optimized, responsive delivery.

■■ This storage implementation requires persistent filesystem storage and is appropriate for the current development/traditional server deployment model. Before deployment to typical serverless infrastructure, uploads should be migrated behind an object-storage adapter (e.g. S3-compatible) — a serverless/edge host's filesystem is typically ephemeral or read-only. No schema change would be needed for that migration; only the save/delete functions in `src/lib/uploads.ts` would change. See [§31](#).

Replacement & removal lifecycle

Both upload routes follow the same strict ordering, so a failure partway through never leaves a user or school without their previous image:

1. The new image is validated and saved to disk first.
2. The database row (`logoUrl/avatarUrl`) is updated.
3. Only once both of those succeed is the *previous* stored file deleted.

A `DELETE` call removes the stored file and clears the field to `null`. Neither route accumulates unused files across normal replace/remove cycles — verified directly (upload → replace → confirm exactly one file remains and the old one is gone; remove → confirm the file is deleted from disk).

Fallback system — `src/components/Avatar.tsx`

One shared component renders both people and schools, with a deliberate visual distinction between the two rather than one generic treatment:

- **People** — a circular avatar. With no photo, a deterministic two-letter initials badge (first + last name initial) in one of the five brand accent colors, chosen by hashing the person's name — the same person always gets the same color, no randomness.
- **Schools** — a softly rounded-square "institution mark," bordered, always a fixed white-background/navy-text treatment (not the person palette) — both to read as distinctly institutional rather than personal, and so the fallback stays legible even placed on a navy background (e.g. the school profile page's hero band), where a randomly-assigned navy fallback would otherwise be invisible.

`/dashboard/profile` — My Profile (My Profile K1)

A self-service page, available to any authenticated role, structured into four sections — identity, institutional relationships, security, and addresses — deliberately not a dashboard (no Homework/Meetings/Calendar/Attendance content lives here):

- **My MEGA Identity** — photo (upload/replace/remove via `ProfilePhotoManager.tsx`), full name, MEGA ID (`user.id`, unchanged) with a client-side **Copy MEGA ID** button (`CopyMegaId.tsx` — `navigator.clipboard.writeText`, falling back to a hidden-textarea `execCommand('copy')`), no server round trip), role chip(s), and email.
- **My Institutional Relationships** (new) — reads `TeacherSchoolAffiliation/StudentSchoolAffiliation/SchoolAdmin` directly (`src/lib/profile.ts`'s `buildInstitutionalRelationships()`), never the legacy `Teacher.schoolId/Student.schoolId` bridge fields and never with a take: 1 — every relationship shows, ACTIVE/PENDING/ENDED alike, including full history. A School Admin's every administered school shows, not one arbitrarily picked. Dates are never fabricated: `startDateSource: "UNKNOWN_MIGRATED"` renders as "Not recorded." Organization relationships are deliberately not shown (see §10a and §37 — `OrganizationAdmin/OrganizationAccountant` have no equivalent status/date/context layer yet).
- **Security & Account** — the existing `ChangePasswordManager.tsx/POST /api/user/password`, unmodified, now under its own heading. Current-password verification, ≥8-character minimum, `bcrypt`, session-scoped authorization, and the demo-account restriction (`isDemoAccountEmail()`, `@megaedu.local` domain, checked server-side) are all unchanged by this kilometer.
- **Addresses** — Current/Permanent, collapsed to a one-line summary (District, Province) by default, expanding into the existing `AddressForm/AddressCard` (byte-identical, untouched) only on Edit, collapsing again after save. Same `PATCH /api/me/address` write path, same Province→District→LocalLevel→Ward cascade, address ownership unchanged (still User, not Student/Teacher).

No new API route was added for the relationships section — it's read server-side, scoped to the session's own `userId`, inside the page component itself. No schema structure change; three stale schema comments claiming this affiliation data was unused ("Phase 1 schema foundation only... nothing reads or writes these yet") were corrected as part of this kilometer. Reachable from a small avatar/"My Profile" link in `SiteHeader.tsx`.

(Source: `src/lib/uploads.ts`, `src/components/Avatar.tsx`, `src/app/api/schools/[id]/logo/route.ts`, `src/app/api/user/avatar/route.ts`, `src/app/dashboard/profile/page.tsx`)

25. Dashboard Architecture

`dashboard/page.tsx` is a single server component that branches on the caller's role-priority order (see §9), resolves the relevant data with direct Prisma queries, and renders the matching role-specific client component (`DashboardClient.tsx`, `TeacherDashboard.tsx`, `StudentDashboard.tsx`, `ParentDashboard.tsx`, `OrgDashboard.tsx`, `PlatformAdminDashboard.tsx`). Each role dashboard's `DashboardHero` now optionally shows an identity badge — the school's logo for the School Admin dashboard, the person's own avatar for Teacher/Student/Parent — via the shared `Avatar` component (see §24). A separate, role-agnostic

`/dashboard/profile` page (also §24) is the self-service surface for managing that avatar, reachable from `SiteHeader.tsx`.

Phase 4D: the `SCHOOL_ADMIN` and `TEACHER` branches resolve institutional context (`getAccessibleSchools()`, §10a) before their existing role-scoped queries — for a person connected to 2+ schools, this is what decides which school's data the dashboard actually shows, independent of the role-priority order that decides *which dashboard* to render. `SchoolChooser.tsx` is the shared client component shown when a choice is needed.

Shared presentational components, reused rather than duplicated per role:

- **AcademicProgressPanel.tsx** — Teaching Progress, Test Results, Recent Attendance, Teacher Evaluations, Assessment Results — parameterized by an explicit audience: `"STUDENT" | "PARENT" | "STAFF"` that filters what's shown (e.g. published-only for Student/Parent, unfiltered for Staff). Rendered once per Student, once per linked child on Parent, and on the Student Profile page.
- **AcademicSnapshot.tsx** — a compact Attendance %/Homework Completion %/Overall Performance summary (N/A, never 0%, when the underlying data is absent), rendered on the Student Profile page and once per linked child on Parent — same three figures, same functions, both places. See [ASSESSMENT AND EVALUATION.md](#).
- **MeetingActions.tsx** — all meeting create/complete/cancel/reschedule/link logic, reused identically on three surfaces (General Evaluations, Subject Evaluations panel, Meetings management page).
- **TodaysHomeworkPanel.tsx** (Phase 1) — renders whatever `fetchTodaysHomework()` returns; rendered once on the Student dashboard and once per linked child on the Parent dashboard. See §17a.

`fetchAcademicProgress()/fetchAssessmentResults()/fetchMeetingsForStudent()` (`src/lib/academicProgress.ts`, `src/lib/assessmentResults.ts`) and `fetchTodaysHomework()` (`src/lib/homework.ts`) are the shared query functions every dashboard branch and the Student Profile page call — never a second, parallel query implementation per surface.

(Source: [ARCHITECTURE.md](#), [ASSESSMENT AND EVALUATION.md](#), [ASSESSMENT_RESULTS.md](#))

26. API Architecture & Major Routes

All routes live under `src/app/api/**/route.ts`, grouped by resource. There is no separate backend — API routes and pages share the same Prisma client and `src/lib` helpers. Response bodies are JSON; a successful response generally includes `{ ok: true, ... }`, an error `{ error: string }`. **No GET list routes exist for any Phase 3 config/results model** — every read happens through the relevant page's own direct Prisma query, a deliberate, consistent convention across every admin/results page in this codebase.

The full, exact inventory (method, path, auth, request/response shape) is maintained in [API.md](#) and is not duplicated here to avoid drift between two copies of the same list. Summary by domain:

Domain	Route prefix	Auth
Auth & registration	<code>/api/auth/*</code>	none / NextAuth
Post-registration affiliation	<code>/api/teacher\ student/{join-school, leave-school, transfer-school}</code> , <code>/api/parent/link-child</code> , <code>/api/schools/create-for-admin</code>	session
Platform Admin	<code>/api/admin/*</code>	<code>requirePlatformAdmin</code>

School directory & profile	<code>/api/schools/[id],</code> <code>.../programs, .../news,</code> <code>.../opportunities, .../logo</code>	<code>requireSchoolAdmin</code>
Staff & students	<code>.../students, .../teachers,</code> <code>.../accountants</code>	<code>requireSchoolAdmin</code>
Sessions & grades (Phase 2)	<code>.../academic-sessions*,</code> <code>.../grades*,</code> <code>.../grade-placements,</code> <code>.../grade-decisions,</code> <code>.../section-assignments</code>	<code>requireSchoolAdmin</code>
Subjects & academic assignment (3A)	<code>.../subjects*,</code> <code>.../teacher-academic-assignments*</code>	<code>requireSchoolAdmin</code>
Operations (3B)	<code>.../class-teacher-assignments*,</code> <code>.../attendance*, .../units*,</code> <code>.../tests*</code>	<code>requireSchoolAdmin OR</code> <code>requireClassTeacher/requireTeacherAssignment</code>
Evaluation & PTM (3C)	<code>.../evaluations*,</code> <code>.../meetings*</code>	<code>requireSchoolAdmin OR</code> <code>requireClassTeacher/requireTeacherAssignment</code>
Homework (Phase 1)	<code>.../homework*</code>	<code>requireTeacherAssignment</code> only (create) / <code>requireSchoolAdmin OR</code> <code>requireTeacherAssignment</code> (edit/publish)
Assessment Framework (3D-1)	<code>.../grading-scales*,</code> <code>.../assessment-frameworks*,</code> <code>.../assessment-framework-assignments*</code>	<code>requireSchoolAdmin</code> only
Assessment Results (3D-2/3/4)	<code>.../components/[id]/results,</code> <code>.../publish,</code> <code>/api/schools/[id]/assessment-results/[resultid]</code>	<code>requireSchoolAdmin OR</code> <code>requireTeacherAssignment</code>
Organizations & courses	<code>/api/organizations/*,</code> <code>/api/courses/*,</code> <code>/api/enrollments/*</code>	<code>requireOrgAdmin/requireCourseOwner/inline</code>
Identity & notifications	<code>/api/interests/*,</code> <code>/api/notifications/*,</code> <code>/api/user/avatar</code>	<code>session/inline</code>
Institutional context (4D)	<code>/api/dashboard/school-context</code>	<code>session</code> (sets the <code>mega_school_ctx</code> preference cookie; never itself a grant — see §10a)

(Source: [API.md](#))

27. Business Rules & Data Integrity Protections

The full, individually-approved rule set — each with its rule, rationale, and applicability — is maintained in [PRODUCT RULES.md](#). It is the tie-breaker whenever a future change seems to conflict with existing behavior: if a rule is documented there, it was a deliberate choice, not an oversight. Categories, summarized:

- **Data integrity & migration discipline** — no Prisma enums; additive-first migrations, verified before cleanup; the SQLite-specific bulk-write transaction pattern (flagged for Postgres rework); legacy fields retired by disuse, never deletion.
- **Snapshot fields vs. live lookups** — freeze display text on anything permanent (Certificate, GradeHistoryAudit); logos are the deliberate live-lookup exception; never hand-type or guess an asset — degrade gracefully to name-only.
- **The "one audited/gatekept write path" pattern** — `issueCourseCertificate()`, `recordGradeDecision()`, `reassignSection()`, `carryForwardEligibleStudents()`, `correctAttendance()`, `updateEvaluationRemarks()`, `correctComponentResult()` — each the sole writer for its concern.
- **New placements are creation, not decisions** — Initial Setup, rollover carry-forward, Add Student's inline placement, and the pending queue's manual placement are all unaudited row *creation*; only a change to an *existing* row is a decision requiring an audit.
- **Never guess when confidence is low** — `matchLegacyGradeText()` returns `null` rather than a low-confidence guess; undecided students are never silently defaulted at rollover.
- **The recurring `NULL ≠ NULL` gap** — documented as a standing reminder: any new model with an optional field in its unique key needs an explicit application-level pre-check, not just the database constraint.
- **Per-phase rules** for Subjects/Teacher Assignment (3A), Academic Operations (3B), Evaluation/PTM (3C), Assessment Framework (3D-1), and Assessment Results/Publishing (3D-2/3/4) — each documented individually in [PRODUCT RULES.md](#) with rule, rationale, and verification evidence.
- **CURRENT_ROSTER_STATUSES** and the **Promotion-panel's deliberately narrower eligibility** — the newest additions, documented in full in [PRODUCT RULES.md](#) and summarized in [§12/§22](#) above.

28. Audit Trails

Model audited	Audit table	Written by	Trigger
GradeHistory (status/outcome)	GradeHistoryAudit	<code>recordGradeDecision()</code>	Every decision, including the first ever made on a row
GradeHistory (section)	GradeHistoryAudit (same table, <code>previousSectionId/newSectionId</code>)	<code>reassignSection()</code>	Every reassignment on an <i>existing</i> row (not the initial set-at-creation value)
Attendance	AttendanceAudit	<code>correctAttendance()</code>	Every correction, capturing both status and remarks even if only one changed
StudentEvaluation	StudentEvaluationAudit	<code>updateEvaluationRemarks()</code>	Only once shared with either Parent or Student — silent while still a private draft

AssessmentComponentResult	AssessmentComponentResult	correctComponentResult()	Only once the parent subject's publication is PUBLISHED — silent while still DRAFT
---------------------------	---------------------------	--------------------------	--

Deliberately *not* audited (current-state, operational, freely re-creatable data, not permanent decisions):
TeacherGradeAssignment, GradeSubject, TeacherAcademicAssignment, ClassTeacherAssignment,
AssessmentFramework/AssessmentPeriod/AssessmentComponent/GradingScale/AssessmentFrameworkAssignment
(structural config, not history), ParentTeacherMeeting.outcomeNotes and rescheduling.

Every audit table in this system is genuinely **append-only** — no route anywhere updates or deletes an audit row.

29. Security Model

- **Session:** JWT-based (NextAuth), not database sessions; `bcrypt.compare` for login.
- **Authorization:** enforced server-side on every write route via the `requireX()` suite (see §8) — never trusted from UI state alone. Verified repeatedly, throughout every phase, via direct `fetch()` calls from an unauthorized session confirming 403, not just that a button was hidden.
- **Data isolation:** every `studentId/schoolId` used in a read is resolved server-side from the caller's own session-derived relationships (a Parent's own `parent.children`, a School Admin's own `SchoolAdmin` row) — never trusted from client-supplied query parameters. Verified live for Parent academic visibility, the Class Overview roster, and the Student Profile page.
- **What's absent ■:** no rate limiting on login/registration/any route; no CSRF protection beyond NextAuth's own; no session revocation ("log out everywhere" — JWTs live until expiry); no OAuth/SSO; no email verification; no password reset flow; no audit log of authorization *decisions* themselves (as distinct from the data-level audit trails in §28).

(Source: [AUTHENTICATION_AND_AUTHORIZATION.md](#), [MEGA_ID.md](#), [KNOWN_GAPS.md](#))

30. Calculation Engines & Shared Libraries

Module	Purpose
<code>auth.ts</code>	NextAuth configuration
<code>authorize.ts</code>	The <code>requireX()</code> guard suite — see §8
<code>prisma.ts</code>	Singleton Prisma client
<code>notify.ts</code>	<code>notify()/notifySchoolCommunity()</code> — best-effort, never allowed to fail the calling action
<code>gradeHistory.ts</code>	<code>recordGradeDecision()</code> , <code>reassignSection()</code> , <code>CURRENT_ROSTER_STATUSES</code>
<code>gradeMatching.ts</code>	<code>matchLegacyGradeText()</code> — never guesses

<code>gradeRollover.ts</code>	<code>carryForwardEligibleStudents()</code> , <code>findPendingStudents()</code> — idempotent, re-runnable
<code>attendance.ts</code>	<code>correctAttendance()</code>
<code>evaluation.ts</code>	<code>updateEvaluationRemarks()</code> , <code>shareEvaluation()</code>
<code>assessmentFramework.ts</code>	<code>resolveFrameworkAssignment()</code> , <code>assignmentCollisionExists()</code> , <code>componentCollisionExists()</code>
<code>assessmentResults.ts</code>	<code>computeComponentContribution()</code> , <code>aggregateGroup()</code> , <code>lookupGrade()</code> , <code>computeUnweightedGPA()</code> , <code>computeUnweightedAveragePercentage()</code> , <code>fetchAssessmentResults()</code> , <code>correctComponentResult()</code> , <code>buildReportCard()</code>
<code>academicProgress.ts</code>	<code>fetchAcademicProgress()</code> , <code>fetchMeetingsForStudent()</code> — the shared, audience-parameterized query functions behind every dashboard branch and the Student Profile page
<code>certificates.ts / certificateView.ts</code>	<code>issueCourseCertificate()</code> , <code>buildCertificateViewModel()</code>

All five "sole write path" functions (`issueCourseCertificate`, `recordGradeDecision`, `reassignSection`, `carryForwardEligibleStudents`, `correctAttendance`) share the same shape: typed input, optional `tx?: Prisma.TransactionClient` for composing into a larger transaction, transactional by default otherwise.

(Source: [ARCHITECTURE.md](#))

31. Deployment & Setup Requirements

Local development

```
npm install npm run dev # next dev — http://localhost:3000
```

Other scripts: `npm run build`, `npm run start`, `npm run lint`, `npm run db:push` (`prisma db push` — direct schema apply, no migration files), `npm run db:seed` (`tsx prisma/seed.ts`, idempotent), `npm run db:studio`.

Database

Development uses SQLite (`prisma/dev.db`). PostgreSQL is the marked, **never-configured-or-tested** production target. Switching requires changing `provider = "sqlite"` to `provider = "postgresql"` in `schema.prisma` and pointing `DATABASE_URL` at a real connection string.

File uploads — School Logos & Profile Photos ■■

Saved to the local filesystem under `public/uploads/` (see [§24](#)). This storage implementation requires persistent filesystem storage and is appropriate for the current development/traditional server deployment model. Before deployment to typical serverless infrastructure, uploads should be migrated behind an object-storage adapter.

Production ■

Nothing has been deployed. No hosting platform, no CI/CD pipeline, no Dockerfile, no production `next.config.js` overrides exist in this repository — genuinely unstarted work, not an oversight in documentation.

Known deployment requirements, inferred from the codebase (not yet acted on)

1. A Postgres database, with the schema provider switched and `npx prisma db push` (or a proper migration) run against it.
2. A real `NEXTAUTH_SECRET/NEXTAUTH_URL` matching the deployed domain.
3. A real `SEED_ADMIN_EMAIL/SEED_ADMIN_PASSWORD` before seeding, or skip seeding in production entirely.
4. **Review the SQLite-specific bulk-write transaction pattern before relying on it under Postgres** (see below) — this is the single most concrete pre-deployment code risk currently documented.
5. Standard Node.js runtime requirements for a Next.js 14 App Router app — no edge-specific code exists anywhere in this codebase.

(Source: [DEPLOYMENT.md](#))

32. Environment Variables

The complete, real list — nothing else is read anywhere in the app:

Variable	Purpose	Dev default
<code>DATABASE_URL</code>	Prisma connection string	<code>file:../dev.db</code>
<code>NEXTAUTH_SECRET</code>	JWT signing secret	placeholder — must be replaced for any real deployment
<code>NEXTAUTH_URL</code>	Canonical app URL for NextAuth callbacks	<code>http://localhost:3000</code>
<code>SEED_ADMIN_EMAIL</code>	Platform Admin account created by <code>db:seed</code>	<code>admin@megaedu.local</code>
<code>SEED_ADMIN_PASSWORD</code>	Same	<code>ChangeMe123!</code> — must be changed before seeding a real environment

(Source: [DEPLOYMENT.md](#))

33. Database Migration & Update Procedures

This project uses `npx prisma db push` — direct schema application, **no `prisma migrate/migrations/` folder** — appropriate for the current single-environment SQLite setup, not necessarily for a multi-environment production deployment.

Established discipline (both Phase 1 and Phase 2 followed the same sequence, enforced manually, not by tooling):

1. Add new models/fields **additively** — nothing existing removed in the same pass.
2. Before applying anything that could conflict with existing data (a new `@@unique` constraint, for example), write a one-off verification script, run it, confirm no conflicts, delete the script.
3. Apply with `npx prisma db push`.
4. Only after a change is applied and verified does any "clean up the old thing" pass happen — and some legacy fields (`Student.gradeLevel`) are intentionally **never** dropped.

■ ■ The Postgres migration caveat, specifically

`grade-placements` and `teacher-assignments` catch a unique-constraint violation (P2002) *inside* an open transaction and continue looping — verified to work correctly on SQLite (a caught statement error doesn't poison the rest of the transaction), but this is **not true on PostgreSQL**, where a failed statement aborts the whole transaction until an explicit rollback. This must be re-checked (likely reworked to pre-filter duplicates before the transaction opens) before any Postgres migration. **Not applicable** to `grade-decisions` or the rollover sweep, which validate eligibility *before* opening their transaction — a genuinely different, already Postgres-safe pattern.

(Source: [ARCHITECTURE.md](#), [PRODUCT RULES.md](#), [DEPLOYMENT.md](#))

34. Development Workflow & Verification

This project has **no automated test suite** (no Jest/Vitest/Playwright/Cypress, no `*.test.*` files anywhere). The established substitute, followed consistently across every phase of this project, is a disciplined manual workflow:

1. **Investigate before designing.** A read-only pass over the relevant schema, `authorize.ts` patterns, and any precedent feature (e.g. reading `certificateView.ts` before designing the Report Card, or `UnitTestResult`'s eager-roster pattern before deciding Assessment Results should be lazy instead) precedes any code.
2. **Design and get explicit approval before implementing.** Non-trivial features (a new phase, a schema change) get a proposed design reviewed and approved before any code is written — the same discipline this very documentation task followed (Table of Contents approved before writing).
3. **Implement additively**, following the migration discipline in [§33](#).
4. **Typecheck as a gate:** `npx tsc --noEmit -p tsconfig.json` after every non-trivial change.
5. **Verify live, against the running dev server and the real database** — not mocked, not assumed. This includes negative cases (confirming an unrelated user is correctly redirected/403'd, not just that the right person succeeds) and, for anything transactional, confirming the transaction actually behaves atomically (a duplicate is skipped without crashing the batch; a partial failure doesn't leave inconsistent rows).
6. **Throwaway verification scripts**, run once via `npx tsx` and then deleted — never left behind in the repository. Used for anything touching real data: checking for existing duplicate rows before adding a constraint, timing a bulk route, proving a sweep is idempotent.
7. **Clean up test-data afterward, with before/after row-count confirmation.** Every phase's verification pass ends by confirming the school's real, pre-existing data (student/teacher/assignment counts) was unaffected by the throwaway fixtures used during testing.
8. **Update documentation in the same pass**, not as an afterthought — if a change alters an architectural or business-rule decision, the relevant `/docs` file is updated (or, for a significant behavior change, fully rewritten) in the same body of work, distinguishing ■/■/■■/■■ status accurately rather than blurring them to make a

feature sound more finished than it is.

9. **Get explicit approval before committing/pushing.** Git history is not touched (`git add/git commit`, let alone `push`) until the user has reviewed the change and explicitly asked for it.

Seeded demo accounts (`prisma/seed.ts`, idempotent) are the de facto manual test fixtures — one account per role, plus a verified demo school/course. See [TESTING.md](#) for the full account list and the verification evidence recorded for every phase.

(Source: [DEVELOPMENT_GUIDELINES.md](#), [TESTING.md](#))

35. Demo & Sample Environment

Beyond the minimal bootstrap seed described above, `prisma/seed-demo.ts` builds an optional, idempotent, self-sufficient demonstration environment on top of it — for testing, demonstration, training, exploring workflows, and multi-school testing.

Two demonstration schools, deliberately different in scale, to demonstrate genuine multi-school data isolation: - **Sunrise Academy** — the flagship, full-depth demonstration environment. - **Himalayan Secondary School** — a smaller, independent school, showing that each school's data (sessions, grades, teachers, students) stays fully isolated from the other.

The environment includes realistic **fictional** examples of schools, teachers, students, parents, academic structures, promotion history, assessment results, attendance, evaluations, unit tests, Parent-Teacher Meetings, and courses/certificates — all generated deterministically (a seeded PRNG), so re-running the seed produces identical results every time.

```
npm run db:seed # baseline bootstrap npm run db:seed:demo # full demo environment npm run db:verify:demo # confirms internal consistency against real calculation logic
```

Full operational detail — the complete account/login list, reset procedure, and what each module contains — is maintained in [DEMO_DATA.md](#) and not duplicated here.

(Source: [DEMO_DATA.md](#))

36. Maintenance Guidelines

For anyone — human or AI-assisted — making a future change to this codebase:

- **Read [PRODUCT_RULES.md](#) before touching business logic.** If a change seems to conflict with a documented rule, that's a signal to stop and check, not to route around it.
- **Read [ARCHITECTURE.md](#) before touching architecture.** Match the established conventions (server components fetch Prisma directly; the "one audited/gatekept write path" pattern; `router.refresh()` as the only revalidation strategy) rather than introducing a new pattern in one corner of the app.
- **Scope discipline:** a bug fix doesn't need surrounding cleanup; a one-shot operation doesn't need a speculative helper. If a request describes 5 steps, build 5 steps.
- **Data integrity:** preserve historical records — no route should delete `GradeHistory`, `GradeHistoryAudit`, or `Certificate` rows. Test migrations for data conflicts before applying a new constraint. Never use a Prisma enum.

- **Don't invent.** If a request assumes something exists that a codebase search doesn't confirm (the documented example: an earlier request assumed "premium courses"/"course bundles" were previously approved — a direct search found zero evidence they were ever built or designed), say so rather than building a plausible-sounding version to satisfy the request's framing.
- **Verify, don't assume.** Typecheck, then exercise the actual change against the running dev server and real database with throwaway fixtures, cleaned up afterward.
- **Update documentation in the same pass** as any architectural or business-rule change, keeping the ■/■/■■/■ distinction honest.

(Source: [DEVELOPMENT_GUIDELINES.md](#))

37. Known Gaps & Deliberate Out-of-Scope Decisions

The complete, individually-re-verified list is maintained in [KNOWN_GAPS.md](#). Organized here by category and status:

■■ Known limitations (real gaps, not by design)

Gap	Detail
School.isActive/Organization.isActive never set to false	Read in two places, no deactivation action exists anywhere
SQLite-specific bulk-write transaction pattern	Would misbehave on PostgreSQL — see §33
No automated test suite	See §34 for the manual substitute
"Roll No." on Class Overview is display-only	Not a persisted schema field — a per-section sequential position computed at render time
School logo / avatar uploads assume a persistent local filesystem	Appropriate for the current dev/traditional-server model; would need an object-storage adapter before serverless deployment — see §24 , §31

■ Deliberately out of scope (approved decisions, not oversights)

Area	Deferred item
Sections	Section-level teacher assignment, section-level analytics/reporting
Subjects/Units	Copying a GradeSubject offering or TeachingPlan/TeachingUnit set forward from a prior session
Teaching structure	Any teaching hierarchy (primary/assistant/substitute teacher)
Unit Tests	Retest concept

Evaluations/PTM	Un-share path for evaluations; parent-initiated meeting requests; meeting recurrence/series
Assessment	Subject-credit/weighting concept (GPA stays unweighted); <code>UnitTest</code> marks auto-derivation; GET list API routes for any Phase 3D model
Promotion	<code>GradingScaleBand.isPassing</code> is not read by any promotion logic — reserved for a future reference display only
Certificates	Grade-certificate issuance (<code>issueGradeCertificate()</code> doesn't exist); PDF export; QR code generation
Commerce	Payment processor integration; premium/bundle course concepts (confirmed via direct search: zero evidence they were ever designed)
Auth	OAuth/SSO, email verification, password reset, rate limiting, session revocation, account deactivation/deletion
Platform Admin	No in-app route to grant/revoke <code>PLATFORM_ADMIN</code>

■ Homework (Phase 1) — deliberate scope decisions

Area	Deferred item
Submissions & grading	No submissions, attachments, grading, rubrics, or feedback — a publish-and-view mechanism only; no per-student result row exists
Notifications	Publishing does not notify anyone (unlike <code>NewsPost</code> 's <code>notifySchoolCommunity()</code>) — a Student/Parent must visit their dashboard
Upcoming/past views	Only "due today" is surfaced; <code>fetchTodaysHomework()</code> 's exact-match <code>dueDate</code> query would need to become a range query
Admin-on-behalf-of-teacher authoring	Unlike <code>Evaluations/TeachingUnit</code> ; <code>Homework.teacherId</code> is a real, non-nullable <code>Teacher</code> FK and Phase 1 has no approved flow for it

■ Calendar (Kilometer 1 / 1.1 / 1.2) — deliberate scope decisions

Area	Deferred item
Nepali calendar	<code>BS↔AD</code> conversion of any kind — no library installed, <code>bsDateDisplay</code> remains unpopulated, no date is computed or fabricated

General Calendar data	Seeded reference dates cover September–December 2026 only, resting on a secondary source not yet confirmed against NPNS/MoHA
Filtering	No Grade/Section/Subject or category filtering on the Admin Calendar
Weekly holiday	Hardcoded to Saturday — no <code>School</code> field, no per-school configuration
Day Status	Overlapping statuses never blend — one dominant value only, by fixed priority
Full Organization Calendar subsystem	Partially resolved (K8) — Organization Event management (create/edit/deactivate, <code>requireOrgAdmin</code> -gated, public display on <code>/organizations/[slug]</code>) now exists; what remains missing is the School-equivalent Annual/Agenda multi-view Calendar page and inclusion in the public <code>/calendar</code> projection layer (<code>src/lib/events.ts</code> untouched by K8)
Interaction model	Recurring events, calendar sync/export/subscriptions, notifications, an interactive Week/Day scheduling grid, drag/drop

■ My Profile (Kilometer 1) — deliberate scope decisions

Area	Deferred item
Identity architecture	No true Person layer — MEGA ID remains <code>User.id</code> ; a <code>Teacher/Student</code> with <code>userId: null</code> has no MEGA ID today
Organization relationships	Not shown — <code>OrganizationAdmin/OrganizationAccountant</code> have no status/date columns or institutional-context layer, the same gap Calendar K1.1 deferred
Account settings	No preferences, privacy, or notification settings — Profile is identity/account/security/relationships/addresses only
Auth features	No Forgot Password, MFA, or session/device management — covered platform-wide under Auth above, not reintroduced here as if newly discovered
School Admin history	<code>SchoolAdmin</code> has no status/startDate/endDate — a row's existence is the entire relationship; no historical join/leave record possible

■ Designed but not implemented

- Grade-based certificates (a reserved field exists on `Certificate`, but no issuance path).

■ Phase 4D — institutional context migration in progress

Gap	Detail
Several dashboard areas still resolve school context via the legacy arbitrary-pick pattern	Initial Setup, New Session, Assessment Frameworks, and Assessment Results still use a plain <code>findFirst({ userId })</code> pick rather than one of the three proven migration patterns (§10a) — in-progress migration debt, not a security gap (every write route independently re-checks ownership of the resource being changed). <code>/dashboard/profile</code> is no longer on this list — My Profile K1 shows every affiliation/administered-school row at once rather than resolving to a single "current" school, so it never needed the chooser pattern in the first place
Organization 2+-organization chooser UX doesn't exist yet	Resolved for the arbitrary-pick itself — <code>dashboard/page.tsx</code> 's <code>ORGANIZATION_ADMIN</code> branch now resolves via <code>getAccessibleOrganizations()</code> (see §10a), never <code>findFirst()</code> . What remains missing is a genuine chooser/switcher UX for 2+ organizations — today that case shows a non-interactive boundary listing, not a working switcher
Student simultaneous multi-school affiliation is an undecided product policy	<code>StudentSchoolAffiliation</code> permits 2+ simultaneous ACTIVE rows (schema-unrestricted, mirroring Teacher), but unlike Teacher this has never been explicitly designed for, tested, or business-approved — nothing blocks it today, nothing was built assuming it happens

(Source: [KNOWN_GAPS.md](#))

38. Appendix: Documentation Index

Document	Covers
PROJECT_OVERVIEW.md	High-level status of every phase, module list, current dev status table
ARCHITECTURE.md	Rendering patterns, <code>src/lib</code> "only path" pattern, module relationships
DATABASE.md	Full field-by-field Prisma model inventory
API.md	Complete API route inventory with request/response shapes
AUTHENTICATION AND AUTHORIZATION.md	Login, session, the full <code>requireX()</code> suite
USER ROLES.md	What every role can do, approval workflows

MEGA ID.md	The identity model and its design principles
INSTITUTIONAL CONTEXT.md	The affiliation lifecycle, <code>getAccessibleSchools()</code> / <code>verifySchoolAccess()</code> , the preference cookie, and the three multi-school migration patterns (Phase 3–4D)
ACADEMIC SESSIONS.md	Sessions, rollover, the one-ACTIVE-per-school rule
GRADES AND PROMOTION.md	Grades, sections, promotion, Class Overview
ACADEMIC STRUCTURE.md	Subjects, <code>GradeSubject</code> , <code>TeacherAcademicAssignment</code> (Phase 3A)
ACADEMIC OPERATIONS.md	Grade Coordinators/Class Teachers, Attendance, Teaching Units, Unit Tests (Phase 3B)
ASSESSMENT AND EVALUATION.md	Evaluations, Parent-Teacher Meetings (Phase 3C)
HOMEWORK.md	Homework model, Draft/Published lifecycle, date semantics, teacher authorization, Student/Parent visibility (Phase 1)
CALENDAR.md	Calendar as a projection layer, Day Status vs Activities, <code>SchoolCalendarEntry</code> , role-aware aggregation, Annual/Agenda views, links (Kilometer 1/1.1/1.2)
ASSESSMENT FRAMEWORK.md	Framework/grading-scale configuration, the guided wizard (Phase 3D-1)
ASSESSMENT RESULTS.md	Marks entry, publishing, calculation engine, Report Card (Phase 3D-2/3/4)
CERTIFICATES.md	Certificate model, issuance, display
COURSES AND ENROLLMENTS.md	MEGA Academy courses and enrollment
PRODUCT RULES.md	Every explicitly-approved business rule, with rationale
KNOWN GAPS.md	Every known gap and deliberate scope decision, individually re-verified
CHANGELOG.md	Dated, hand-written record of every notable change
DEPLOYMENT.md	Local dev setup, environment variables, production readiness
DEVELOPMENT GUIDELINES.md	Rules for future work on this codebase
TESTING.md	Verification practice, seeded demo accounts
DEMO DATA.md	The full demo/sample environment: schools, accounts, login info, seed/reset/verify commands

Note on the generated PDF snapshots: docs/MEGA_EDU_Technical_Documentation_YYYY-MM-DD.pdf (dated snapshots exist for 2026-08-31, 2026-09-01, and 2026-09-05) are generated exports of this Markdown document at a point in time — this Markdown file is the source and the current technical reference; a dated PDF reflects whatever this file said as of its own date and is not updated retroactively. The two most recently dated snapshots (2026-08-31 and, per current protection policy, any other already-generated dated snapshot) are treated as protected artifacts and must never be regenerated, modified, or replaced as a side effect of editing this source file — see [DEVELOPMENT GUIDELINES.md](#). As of this edit, no dated snapshot yet reflects Calendar (Kilometer 1/1.1/1.2) or My Profile (Kilometer 1) — regenerating a new dated snapshot from this file is a separate, explicitly-authorized action, not implied by updating the source.